

1. Download the data from [here](#). You have to use data.csv file for this assignment
2. Code the model to classify data like below image. You can use any number of units in your Dense layers.

3. Writing Callbacks

You have to implement the following callbacks

- Write your own callback function, that has to print the micro F1 score and AUC score after each epoch. Do not use tf.keras.metrics for calculating AUC and F1 score.
- Save your model at every epoch if your validation accuracy is improved from previous epoch.
- You have to decay learning based on below conditions Cond1. If your validation accuracy at that epoch is less than previous epoch accuracy, you have to decrease the learning rate by 10%. Cond2. For every 3rd epoch, decay your learning rate by 5%.
- If you are getting any NaN values (either weights or loss) while training, you have to terminate your training.
- You have to stop the training if your validation accuracy is not increased in last 2 epochs.
- Use tensorboard for every model and analyse your scalar plots and histograms. (you need to upload the screenshots and write the observations for each model for evaluation)

Note

Make sure that you are plotting tensorboard plots either in your notebook or you can try to create a pdf file with all the tensorboard screenshots. Please write your analysis of tensorboard results for each model.

Import Section

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
import os
import datetime
from sklearn.metrics import f1_score, roc_auc_score
from sklearn.model_selection import train_test_split
```

```

from sklearn.preprocessing import StandardScaler
from tensorflow.keras.callbacks import
ModelCheckpoint,EarlyStopping,LearningRateScheduler,ReduceLRonPlateau,
TensorBoard
from tensorflow.keras.layers import Dense,Input,Activation,Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.utils import plot_model
from collections import Iterable
import warnings
warnings.filterwarnings('ignore')

```

Lets download the dataset

```
!gdown --id 15dCNcmKskcFVjs7R0ElQkR61Ex53uJpM
```

```

/usr/local/lib/python3.8/dist-packages/gdown/cli.py:127:
FutureWarning: Option `--id` was deprecated in version 4.3.1 and will
be removed in 5.0. You don't need to pass it anymore to use a file ID.
  warnings.warn(
Downloading...
From: https://drive.google.com/uc?id=15dCNcmKskcFVjs7R0ElQkR61Ex53uJpM
To: /content/data.csv
100% 887k/887k [00:00<00:00, 83.2MB/s]

```

```

dataset = pd.read_csv('/content/data.csv')
dataset.head()

```

```

      f1      f2  label
0  0.450564  1.074305    0.0
1  0.085632  0.967682    0.0
2  0.117326  0.971521    1.0
3  0.982179 -0.380408    0.0
4 -0.720352  0.955850    0.0

```

```
dataset.describe()
```

```

      f1      f2      label
count  20000.000000  20000.000000  20000.000000
mean      0.000630   -0.000745    0.500000
std      0.671165    0.674704    0.500013
min     -1.649781   -1.600645    0.000000
25%     -0.589878   -0.596424    0.000000
50%      0.001795   -0.003113    0.500000
75%      0.586631    0.597803    1.000000
max      1.629722    1.584291    1.000000

```

```
dataset['label'].value_counts()
```


```

0.0    10000
1.0    10000
Name: label, dtype: int64


```

2-D Scatter Plot

```
sns.scatterplot(data=dataset, x="f1", y="f2", hue="label")  
<matplotlib.axes._subplots.AxesSubplot at 0x7f33e40cb400>
```



```
sns.pairplot(dataset, hue="label")  
<seaborn.axisgrid.PairGrid at 0x7f0aeb58b310>
```



Lets create train and test datasets

```
X = dataset.iloc[:, :-1]
y = dataset.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.30, stratify = y, random_state = 42)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(14000, 2)
(6000, 2)
(14000,)
(6000,)
```

Write your own callback function, that has to print the micro F1 score and AUC score after each epoch. Do not use tf.keras.metrics for calculating AUC and F1 score.

```
class Custom_AUC_f1(tf.keras.callbacks.Callback):
    def __init__(self, validation_data):
        self.x_test = validation_data[0]
        self.y_test = validation_data[1]
    def on_train_begin(self, logs={}):
        self.val_f1s = []
        self.val_rocs = []
    def on_epoch_end(self, epoch, logs={}):
        y_predict = (np.asarray(self.model.predict(self.x_test))).round()
        val_f1 = f1_score(y_test, y_predict.round())
        roc_val = roc_auc_score(y_test, y_predict)
        self.val_f1s.append(val_f1)
        self.val_rocs.append(roc_val)
        print("-f1 score :", val_f1, "-ROCValue :", roc_val)
```

```
Custom_AUC_F1 = Custom_AUC_f1(validation_data=[X_test, y_test])
```

Save your model at every epoch if your validation accuracy is improved from previous epoch.

```
filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath,
                             monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')
```

You have to decay learning based on below conditions

Cond1: If your validation accuracy at that epoch is less than previous epoch accuracy, you have to decrease the learning rate by 10%.

Cond2: For every 3rd epoch, decay your learning rate by 5%.

```
reduce_lr = ReduceLROnPlateau(monitor='val_accuracy',
                              factor=0.9, patience=1, min_lr=0.0001)
```

```
def changeLearningRate(epoch, lr):
    #here we are performing exponential decay of the learning rate
    if (epoch+1)% 3 == 0:
        return lr*0.95
    return lr
lrschedule = LearningRateScheduler(changeLearningRate, verbose=1)
```

If you are getting any NaN values (either weights or loss) while training, you have to terminate your training.

```
class TerminateNaN(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        loss = logs.get('loss')
```

```

        if loss is not None:
            if np.isnan(loss) or np.isinf(loss):
                print("Invalid loss and terminated at epoch
{}".format(epoch))
                self.model.stop_training = True
        def flatten(lis):
            for item in lis:
                if isinstance(item, Iterable) and not isinstance(item,
str):
                    for x in flatten(item):
                        yield x
                    else:
                        yield item
        weights =
np.array(list(flatten(np.array(self.model.get_weights()))))
        if(np.isnan(weights).any() or np.isinf(weights).any()):
            print("Invalid weights and terminated at epoch
{}".format(epoch))
            self.model.stop_training = True

terminate = TerminateNaN()

```

You have to stop the training if your validation accuracy is not increased in last 2 epochs.

```

earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.35,
patience=2, verbose=1)

```

Use tensorboard for every model and analyse your scalar plots and histograms.

```

# Load the TensorBoard notebook extension
%load_ext tensorboard
log_dir = os.path.join("logs", 'fits',
datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback =
tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_
graph=True)
%reload_ext tensorboard

```

Model-1

1. Use tanh as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
4. Analyze your output and training process.

```

X_train.shape

```

```

(14000, 2)

```

```

#Input layer

```

```

input_layer = Input(shape=(X_train.shape[1],))

```

```

#Dense hidden layer

```

```

layer1 =
Dense(10,activation='tanh',kernel_initializer=tf.keras.initializers.Ra
ndomUniform(minval=-0, maxval=1))(input_layer)
#Dense hidden layer
layer2 =
Dense(20,activation='tanh',kernel_initializer=tf.keras.initializers.Ra
ndomUniform(minval=-0, maxval=1))(layer1)
#Dense hidden layer
layer3 =
Dense(30,activation='tanh',kernel_initializer=tf.keras.initializers.Ra
ndomUniform(minval=-0, maxval=1))(layer2)
#Dense hidden layer
layer4 =
Dense(20,activation='tanh',kernel_initializer=tf.keras.initializers.Ra
ndomUniform(minval=-0, maxval=1))(layer3)
#Dense hidden layer
layer5 =
Dense(10,activation='tanh',kernel_initializer=tf.keras.initializers.Ra
ndomUniform(minval=-0, maxval=1))(layer4)
#output layer
output =
Dense(1,activation='sigmoid',kernel_initializer=tf.keras.initializers.
RandomUniform(minval=-0, maxval=1))(layer5)
#Creating a model
model_one = Model(inputs=input_layer,outputs=output)
model_one.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 2)]	0
dense (Dense)	(None, 10)	30
dense_1 (Dense)	(None, 20)	220
dense_2 (Dense)	(None, 30)	630
dense_3 (Dense)	(None, 20)	620
dense_4 (Dense)	(None, 10)	210
dense_5 (Dense)	(None, 1)	11
=====		
Total params: 1,721		
Trainable params: 1,721		
Non-trainable params: 0		
=====		

```
plot_model(model_one, show_shapes=True)
```


input_1	input:	[(None, 2)]
InputLayer	output:	[(None, 2)]



dense	input:	(None, 2)
Dense	output:	(None, 10)



dense_1	input:	(None, 10)
Dense	output:	(None, 20)



dense_2	input:	(None, 20)
Dense	output:	(None, 30)



dense_3	input:	(None, 30)
Dense	output:	(None, 20)



dense_4	input:	(None, 20)
Dense	output:	(None, 10)



#Callbacks

```
model_one.compile(optimizer='sgd',  
loss='binary_crossentropy',metrics=['accuracy'])  
model_one.fit(X_train,y_train,epochs=10,  
validation_data=(X_test,y_test), batch_size=16,  
callbacks=[Custom_AUC_F1,checkpoint,reduce_lr,lrschedule,terminate,earlystop,tensorboard_callback])
```

Epoch 1: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 1/10

1/875 [.....] - ETA: 15:21 - loss: 2.2287 - accuracy: 0.5625

WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0024s vs `on_train_batch_end` time: 0.0026s). Check your callbacks.

188/188 [=====] - 0s 1ms/step
-f1 score : 0.49246401354784086 -ROCValue : 0.5005

Epoch 1: val_accuracy improved from -inf to 0.50050, saving model to model_save/weights-01-0.5005.hdf5

875/875 [=====] - 5s 4ms/step - loss: 0.7930
- accuracy: 0.4991 - val_loss: 0.6932 - val_accuracy: 0.5005 - lr: 0.0100

Epoch 2: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 2/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.5076442544796975 -ROCValue : 0.5008333333333333

Epoch 2: val_accuracy improved from 0.50050 to 0.50083, saving model to model_save/weights-02-0.5008.hdf5

875/875 [=====] - 3s 3ms/step - loss: 0.6936
- accuracy: 0.5060 - val_loss: 0.6931 - val_accuracy: 0.5008 - lr: 0.0100

Epoch 3: LearningRateScheduler setting learning rate to 0.009499999787658453.

Epoch 3/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.459092533047517 -ROCValue : 0.4953333333333333

Epoch 3: val_accuracy did not improve from 0.50083

875/875 [=====] - 3s 3ms/step - loss: 0.6936

- accuracy: 0.5016 - val_loss: 0.6932 - val_accuracy: 0.4953 - lr: 0.0085

Epoch 3: early stopping

<keras.callbacks.History at 0x7f0ae30a1d60>

%tensorboard --logdir logs

Output hidden; open in <https://colab.research.google.com> to view.

Model-2

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
4. Analyze your output and training process.

#Input layer

input_layer = Input(shape=(X_train.shape[1],))

#Dense hidden layer

layer1 =

Dense(10,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(input_layer)

#Dense hidden layer

layer2 =

Dense(20,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(layer1)

#Dense hidden layer

layer3 =

Dense(30,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(layer2)

#Dense hidden layer

layer4 =

Dense(20,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(layer3)

#Dense hidden layer

layer5 =

Dense(10,activation='relu',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(layer4)

#output layer

output =

Dense(1,activation='sigmoid',kernel_initializer=tf.keras.initializers.RandomUniform(minval=-0, maxval=1))(layer5)

#Creating a model

model_two = Model(inputs=input_layer,outputs=output)

model_two.summary()

Model: "model_1"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

=====		
input_2 (InputLayer)	[(None, 2)]	0
dense_6 (Dense)	(None, 10)	30
dense_7 (Dense)	(None, 20)	220
dense_8 (Dense)	(None, 30)	630
dense_9 (Dense)	(None, 20)	620
dense_10 (Dense)	(None, 10)	210
dense_11 (Dense)	(None, 1)	11

=====

Total params: 1,721
Trainable params: 1,721
Non-trainable params: 0

plot_model(model_two, show_shapes=True)

input_2	input:	[(None, 2)]
InputLayer	output:	[(None, 2)]



dense_6	input:	(None, 2)
Dense	output:	(None, 10)



dense_7	input:	(None, 10)
Dense	output:	(None, 20)



dense_8	input:	(None, 20)
Dense	output:	(None, 30)



dense_9	input:	(None, 30)
Dense	output:	(None, 20)



dense_10	input:	(None, 20)
Dense	output:	(None, 10)



#Callbacks

```
optimizers_sgd_with_momemtum =  
tf.keras.optimizers.SGD(learning_rate=0.01,momentum=0.7,nesterov=True)  
  
model_two.compile(optimizer=optimizers_sgd_with_momemtum,  
loss='binary_crossentropy',metrics=['accuracy'])  
model_two.fit(X_train,y_train,epochs=10,  
validation_data=(X_test,y_test), batch_size=16,  
callbacks=[Custom_AUC_F1,checkpoint,reduce_lr,lrschedule,terminate,ear  
lystop,tensorboard_callback])
```

Epoch 1: LearningRateScheduler setting learning rate to
0.009999999776482582.

Epoch 1/10

1/875 [.....] - ETA: 8:15 - loss: 4027.7900
- accuracy: 0.2500

WARNING:tensorflow:Callback method `on_train_batch_end` is slow
compared to the batch time (batch time: 0.0023s vs
`on_train_batch_end` time: 0.0027s). Check your callbacks.

188/188 [=====] - 0s 1ms/step
-f1 score : 0.0 -ROCValue : 0.5

Epoch 1: val_accuracy did not improve from 0.50083

875/875 [=====] - 4s 4ms/step - loss: 5.2957
- accuracy: 0.4976 - val_loss: 0.6932 - val_accuracy: 0.5000 - lr:
0.0100

Epoch 2: LearningRateScheduler setting learning rate to
0.009999999776482582.

Epoch 2/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.6666666666666666 -ROCValue : 0.5

Epoch 2: val_accuracy did not improve from 0.50083

875/875 [=====] - 3s 3ms/step - loss: 0.6933
- accuracy: 0.4973 - val_loss: 0.6932 - val_accuracy: 0.5000 - lr:
0.0090

Epoch 3: LearningRateScheduler setting learning rate to
0.008549999631941318.

Epoch 3/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.6666666666666666 -ROCValue : 0.5

Epoch 3: val_accuracy did not improve from 0.50083

```
875/875 [=====] - 3s 3ms/step - loss: 0.6933  
- accuracy: 0.4923 - val_loss: 0.6933 - val_accuracy: 0.5000 - lr:  
0.0077
```

Epoch 3: early stopping

<keras.callbacks.History at 0x7f0adf04d2e0>

%tensorboard --logdir logs

Output hidden; open in <https://colab.research.google.com> to view.

Model-3

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he_uniform() as initializer.
4. Analyze your output and training process.

#Input layer

```
input_layer = Input(shape=(X_train.shape[1],))
```

#Dense hidden layer

```
layer1 =
```

```
Dense(10,activation='relu',kernel_initializer=tf.keras.initializers.he  
_uniform())(input_layer)
```

#Dense hidden layer

```
layer2 =
```

```
Dense(20,activation='relu',kernel_initializer=tf.keras.initializers.he  
_uniform())(layer1)
```

#Dense hidden layer

```
layer3 =
```

```
Dense(30,activation='relu',kernel_initializer=tf.keras.initializers.he  
_uniform())(layer2)
```

#Dense hidden layer

```
layer4 =
```

```
Dense(20,activation='relu',kernel_initializer=tf.keras.initializers.he  
_uniform())(layer3)
```

#Dense hidden layer

```
layer5 =
```

```
Dense(10,activation='relu',kernel_initializer=tf.keras.initializers.he  
_uniform())(layer4)
```

#output layer

```
output =
```

```
Dense(1,activation='sigmoid',kernel_initializer=tf.keras.initializers.  
he_uniform())(layer5)
```

#Creating a model

```
model_three = Model(inputs=input_layer,outputs=output)
```

```
model_three.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
=====		

input_3 (InputLayer)	[(None, 2)]	0
dense_12 (Dense)	(None, 10)	30
dense_13 (Dense)	(None, 20)	220
dense_14 (Dense)	(None, 30)	630
dense_15 (Dense)	(None, 20)	620
dense_16 (Dense)	(None, 10)	210
dense_17 (Dense)	(None, 1)	11

```
=====
Total params: 1,721
Trainable params: 1,721
Non-trainable params: 0
```

```
plot_model(model_three, show_shapes=True)
```


input_3	input:	[(None, 2)]
InputLayer	output:	[(None, 2)]



dense_12	input:	(None, 2)
Dense	output:	(None, 10)



dense_13	input:	(None, 10)
Dense	output:	(None, 20)



dense_14	input:	(None, 20)
Dense	output:	(None, 30)



dense_15	input:	(None, 30)
Dense	output:	(None, 20)



dense_16	input:	(None, 20)
Dense	output:	(None, 10)



#Callbacks

```
optimizers_sgd_with_momemtum =  
tf.keras.optimizers.SGD(learning_rate=0.01,momentum=0.5,nesterov=True)
```

```
model_three.compile(optimizer=optimizers_sgd_with_momemtum,  
loss='binary_crossentropy',metrics=['accuracy'])  
model_three.fit(X_train,y_train,epochs=10,  
validation_data=(X_test,y_test), batch_size=16,  
callbacks=[Custom_AUC_F1,checkpoint,reduce_lr,lrschedule,terminate,earl  
lystop,tensorboard_callback])
```

Epoch 1: LearningRateScheduler setting learning rate to
0.009999999776482582.

Epoch 1/10

1/875 [.....] - ETA: 17:47 - loss: 0.7128 -
accuracy: 0.3750

WARNING:tensorflow:Callback method `on_train_batch_end` is slow
compared to the batch time (batch time: 0.0031s vs
`on_train_batch_end` time: 0.0137s). Check your callbacks.

188/188 [=====] - 0s 2ms/step
-f1 score : 0.6203506472226774 -ROCValue : 0.6138333333333333

Epoch 1: val_accuracy improved from 0.50083 to 0.61383, saving model
to model_save/weights-01-0.6138.hdf5

875/875 [=====] - 6s 5ms/step - loss: 0.6717
- accuracy: 0.5809 - val_loss: 0.6506 - val_accuracy: 0.6138 - lr:
0.0100

Epoch 2: LearningRateScheduler setting learning rate to
0.009999999776482582.

Epoch 2/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.6764005713378829 -ROCValue : 0.6601666666666666

Epoch 2: val_accuracy improved from 0.61383 to 0.66017, saving model
to model_save/weights-02-0.6602.hdf5

875/875 [=====] - 4s 4ms/step - loss: 0.6248
- accuracy: 0.6481 - val_loss: 0.6152 - val_accuracy: 0.6602 - lr:
0.0100

Epoch 3: LearningRateScheduler setting learning rate to
0.009499999787658453.

Epoch 3/10

188/188 [=====] - 0s 1ms/step
-f1 score : 0.5922848664688427 -ROCValue : 0.6565000000000001

Epoch 3: val_accuracy did not improve from 0.66017

```
875/875 [=====] - 4s 5ms/step - loss: 0.6068  
- accuracy: 0.6639 - val_loss: 0.6226 - val_accuracy: 0.6565 - lr:  
0.0085
```

Epoch 3: early stopping

```
<keras.callbacks.History at 0x7f0ae9d93400>
```

```
%tensorboard --logdir logs
```

Output hidden; open in <https://colab.research.google.com> to view.

Model-4

1. Try with any values to get better accuracy/f1 score.

```
!rm -rf ./logs/
```

```
tf.compat.v1.reset_default_graph()
```

Lets employ few techniques to better the model

Lets resplit the data with 90:10 as deep learning perform better with more data.

```
X_model4 = dataset.iloc[:, :-1]  
y_model4 = dataset.iloc[:, -1]  
X_train_model4, X_test_model4, y_train_model4, y_test_model4 =  
train_test_split(X_model4, y_model4, test_size=0.10, stratify =  
y_model4 , random_state = 42)  
print(X_train_model4.shape)  
print(X_test_model4.shape)  
print(y_train_model4.shape)  
print(y_test_model4.shape)
```

```
(18000, 2)
```

```
(2000, 2)
```

```
(18000, )
```

```
(2000, )
```

Lets rescale the data using minmaxscaler and bring data in range 0-1.

```
X_train_model4.describe()
```

	f1	f2
count	18000.000000	18000.000000
mean	0.000661	-0.003546
std	0.670208	0.674711
min	-1.649781	-1.600645
25%	-0.588755	-0.598939
50%	0.000253	-0.008212
75%	0.586433	0.593109
max	1.629722	1.584291

```
from sklearn.preprocessing import MinMaxScaler  
minmaxscaler = MinMaxScaler()
```

```

X_train_model4 = minmaxscalar.fit_transform(X_train_model4)
X_test_model4 = minmaxscalar.transform(X_test_model4)

print(X_train_model4[:,0].min())
print(X_train_model4[:,0].max())
print(X_test_model4[:,0].min())
print(X_test_model4[:,0].max())

0.0
1.0
0.01485105225468586
0.9655135024755643

```

- We can see values are in range 0 to 1

Lets find the perfect hyper parameters using keras tuner

#pip install keras-tuner

```

import keras_tuner as kt
from tensorflow import keras

def build_model(hp):
    model = keras.Sequential()
    counter = 0
    for i in range(hp.Int('num_layers',min_value = 1, max_value = 5)):
        if counter ==0:
            model.add(
                keras.layers.Dense(
                    hp.Int('units'+str(i), min_value = 8, max_value =
100,step =8),
                    activation =
hp.Choice('activation'+str(i),values=['relu','tanh','sigmoid']),
                    kernel_initializer =
hp.Choice('initialiser'+str(i),values =
['he_normal','glorot_normal']),
                    input_dim = X_train_model4.shape[1]

                )
            )
            model.add(
keras.layers.Dropout(hp.Choice('dropout'+str(i),values=[0.1,0.2,0.3,0.
4,0.5,0.6,0.7,0.8,0.9])))
        else:
            model.add(
                keras.layers.Dense(
                    hp.Int('units'+str(i), min_value = 8, max_value =
100,step =8),
                    activation =
hp.Choice('activation'+str(i),values=['relu','tanh','sigmoid']),

```

```

        kernel_initializer =
hp.Choice('initialiser'+str(i), values = ['he_normal', 'glorot_normal'])
    )
    )
    model.add(

keras.layers.Dropout(hp.Choice('dropout'+str(i), values=[0.1,0.2,0.3,0.
4,0.5,0.6,0.7,0.8,0.9])))
    )
    counter += 1
    model.add(keras.layers.Dense(1, activation = 'sigmoid'))
    model.compile(optimizer = hp.Choice('optimiser', values =
['rmsprop', 'sgd', 'adam', 'adadelta', 'nadam']), loss='binary_crossentropy',
metrics=['accuracy'])
    return model

```

```

tuner = kt.RandomSearch(build_model, objective =
'val_accuracy', max_trials = 3, directory =
"hyperparametertuning", project_name = 'callbackassignment1')

```

```

tuner.search(X_train_model4, y_train_model4, epochs = 5, validation_data
= (X_test_model4, y_test_model4))

```

Trial 3 Complete [00h 00m 11s]
val_accuracy: 0.5

Best val_accuracy So Far: 0.6704999804496765
Total elapsed time: 00h 00m 28s

lets get the best hyper parameters

```

tuner.get_best_hyperparameters()[0].values

```

```

{'num_layers': 2,
 'units0': 96,
 'activation0': 'relu',
 'initialiser0': 'glorot_normal',
 'dropout0': 0.4,
 'optimiser': 'adam',
 'units1': 8,
 'activation1': 'relu',
 'initialiser1': 'he_normal',
 'dropout1': 0.1}

```

We got the best hyypeameters lets build the model 4

#Input layer

```

input_layer = Input(shape=(X_train_model4.shape[1],))

```

#Dense hidden layer

```

layer1 =

```

```

Dense(96, activation='relu', kernel_initializer='glorot_normal')

```

```

(input_layer)
#Dropout layer
dropout1 = Dropout(0.4)(layer1)
#Dense hidden layer
layer2 = Dense(8,activation='relu',kernel_initializer='he_normal')(dropout1)
#Dropout layer
dropout2 = Dropout(0.1)(layer2)
#output layer
output = Dense(1,activation='sigmoid')(dropout2)
#Creating a model
model_four = Model(inputs=input_layer,outputs=output)
model_four.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 2)]	0
dense_4 (Dense)	(None, 96)	288
dropout_3 (Dropout)	(None, 96)	0
dense_5 (Dense)	(None, 8)	776
dropout_4 (Dropout)	(None, 8)	0
dense_6 (Dense)	(None, 1)	9

```

Total params: 1,073
Trainable params: 1,073
Non-trainable params: 0

```

```

plot_model(model_four, show_shapes=True)

```

input_1	input:	[(None, 2)]
InputLayer	output:	[(None, 2)]



dense_4	input:	(None, 2)
Dense	output:	(None, 96)



dropout_3	input:	(None, 96)
Dropout	output:	(None, 96)



dense_5	input:	(None, 96)
Dense	output:	(None, 8)



dropout_4	input:	(None, 8)
Dropout	output:	(None, 8)



dense_6	input:	(None, 8)
Dense	output:	(None, 1)

```

reduce_lr_model4 = ReduceLROnPlateau(monitor='val_accuracy',
factor=0.9,patience=1, min_lr=0.0001)

earlystop_model4 = EarlyStopping(monitor='val_accuracy',
min_delta=0.00009, patience=3, verbose=1)

#Callbacks
model_four.compile(optimizer='adam',
loss='binary_crossentropy',metrics=['accuracy'])
model_four.fit(X_train_model4,y_train_model4,epochs=100,
validation_data=(X_test_model4,y_test_model4), batch_size=512,
callbacks=[Custom_AUC_F1,checkpoint,reduce_lr_model4,earlystop_model4,
tensorboard_callback])

```

Epoch 1/100

```

188/188 [=====] - 0s 1ms/step
-f1 score : 0.13624141021810576 -ROCValue : 0.5181666666666667

```

Epoch 1: val_accuracy did not improve from 0.69600

```

36/36 [=====] - 1s 19ms/step - loss: 0.6176 -
accuracy: 0.6606 - val_loss: 0.5937 - val_accuracy: 0.6915 - lr:
0.0010

```

Epoch 2/100

```

188/188 [=====] - 0s 1ms/step
-f1 score : 0.13727245598328858 -ROCValue : 0.5181666666666667

```

Epoch 2: val_accuracy did not improve from 0.69600

```

36/36 [=====] - 1s 15ms/step - loss: 0.6163 -
accuracy: 0.6614 - val_loss: 0.5929 - val_accuracy: 0.6955 - lr:
0.0010

```

Epoch 3/100

```

188/188 [=====] - 0s 1ms/step
-f1 score : 0.13825983313468415 -ROCValue : 0.5179999999999999

```

Epoch 3: val_accuracy improved from 0.69600 to 0.69750, saving model to model_save/weights-03-0.6975.hdf5

```

36/36 [=====] - 1s 15ms/step - loss: 0.6140 -
accuracy: 0.6644 - val_loss: 0.5929 - val_accuracy: 0.6975 - lr:
0.0010

```

Epoch 4/100

```

188/188 [=====] - 0s 1ms/step
-f1 score : 0.13714967203339296 -ROCValue : 0.5176666666666667

```

Epoch 4: val_accuracy did not improve from 0.69750

```

36/36 [=====] - 0s 13ms/step - loss: 0.6170 -
accuracy: 0.6604 - val_loss: 0.5934 - val_accuracy: 0.6940 - lr:
0.0010

```

Epoch 5/100

```

188/188 [=====] - 0s 1ms/step
-f1 score : 0.13714967203339296 -ROCValue : 0.5176666666666667

```


Epoch 5: val_accuracy did not improve from 0.69750
36/36 [=====] - 1s 15ms/step - loss: 0.6150 -
accuracy: 0.6632 - val_loss: 0.5936 - val_accuracy: 0.6935 - lr:
0.0010
Epoch 6/100
188/188 [=====] - 0s 1ms/step
-f1 score : 0.1389385807990459 -ROCValue : 0.5186666666666666

Epoch 6: val_accuracy improved from 0.69750 to 0.69800, saving model
to model_save/weights-06-0.6980.hdf5
36/36 [=====] - 1s 15ms/step - loss: 0.6146 -
accuracy: 0.6647 - val_loss: 0.5929 - val_accuracy: 0.6980 - lr:
0.0010
Epoch 7/100
188/188 [=====] - 0s 1ms/step
-f1 score : 0.14000595770032767 -ROCValue : 0.5188333333333334

Epoch 7: val_accuracy improved from 0.69800 to 0.70050, saving model
to model_save/weights-07-0.7005.hdf5
36/36 [=====] - 1s 15ms/step - loss: 0.6161 -
accuracy: 0.6610 - val_loss: 0.5924 - val_accuracy: 0.7005 - lr:
0.0010
Epoch 8/100
188/188 [=====] - 0s 1ms/step
-f1 score : 0.13825983313468415 -ROCValue : 0.5179999999999999

Epoch 8: val_accuracy did not improve from 0.70050
36/36 [=====] - 0s 13ms/step - loss: 0.6158 -
accuracy: 0.6618 - val_loss: 0.5926 - val_accuracy: 0.6960 - lr:
0.0010
Epoch 9/100
188/188 [=====] - 0s 1ms/step
-f1 score : 0.14051801131289074 -ROCValue : 0.5188333333333334

Epoch 9: val_accuracy did not improve from 0.70050
36/36 [=====] - 0s 13ms/step - loss: 0.6147 -
accuracy: 0.6645 - val_loss: 0.5918 - val_accuracy: 0.6990 - lr:
0.0010
Epoch 10/100
188/188 [=====] - 0s 1ms/step
-f1 score : 0.1394517282479142 -ROCValue : 0.5186666666666667

Epoch 10: val_accuracy did not improve from 0.70050
36/36 [=====] - 0s 13ms/step - loss: 0.6137 -
accuracy: 0.6647 - val_loss: 0.5930 - val_accuracy: 0.6940 - lr:
0.0010
Epoch 10: early stopping
<keras.callbacks.History at 0x7f33da1407c0>

```
%tensorboard --logdir logs
```

Output hidden; open in <https://colab.research.google.com> to view.

Observations:

Accuracy of model 1: 0.49234

Accuracy of model 2: 0.50083

Accuracy of model 3: 0.66017

Accuracy of model 4: 0.70050

Note: Successfully improved the model performance using hyper parameter tuning.