

CNN on CIFR Assignment:

1. Please visit this link to access the state-of-art DenseNet code for reference - [DenseNet - cifar10 notebook link](#)
2. You need to create a copy of this and "retrain" this model to achieve 90+ test accuracy.
3. You cannot use DropOut layers.
4. You MUST use Image Augmentation Techniques.
5. You cannot use an already trained model as a beginning points, you have to initilize as your own
6. You cannot run the program for more than 300 Epochs, and it should be clear from your log, that you have only used 300 Epochs
7. You cannot use test images for training the model.
8. You cannot change the general architecture of DenseNet (which means you must use Dense Block, Transition and Output blocks as mentioned in the code)
9. You are free to change Convolution types (e.g. from 3x3 normal convolution to Depthwise Separable, etc)
10. You cannot have more than 1 Million parameters in total
11. You are free to move the code from Keras to Tensorflow, Pytorch, MXNET etc.
12. You can use any optimization algorithm you need.
13. You can checkpoint your model and retrain the model from that checkpoint so that no need of training the model from first if you lost at any epoch while training. You can directly load that model and Train from that epoch.

Import Section

```
#!/pip install tensorflow-gpu==2.8.3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from tqdm import tqdm
tqdm.pandas()
import tensorflow as tf
import math
import timeit
from six.moves import cPickle as pickle
import os
import datetime
import platform
from subprocess import check_output
from tensorflow.keras import models, layers
from tensorflow.keras.models import Model
from tensorflow.keras.layers import BatchNormalization, Activation,
Flatten,MaxPooling2D
```

```

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import
ModelCheckpoint,EarlyStopping,LearningRateScheduler,ReduceLRonPlateau,
TensorBoard
from tensorflow.keras.utils import plot_model
from keras.preprocessing.image import ImageDataGenerator

import tensorflow
tensorflow.__version__

{"type":"string"}

```

Lets pull the CIFAR 10 dataset

Load CIFAR10 Data

```

(X_train, y_train), (X_test, y_test) =
tf.keras.datasets.cifar10.load_data()
img_height, img_width, channel =
X_train.shape[1],X_train.shape[2],X_train.shape[3]

```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 2s 0us/step

Lets Analyse the data

- Shape of each image is 32X32X3 that means it is a color image

```

X_train[0].shape

(32, 32, 3)

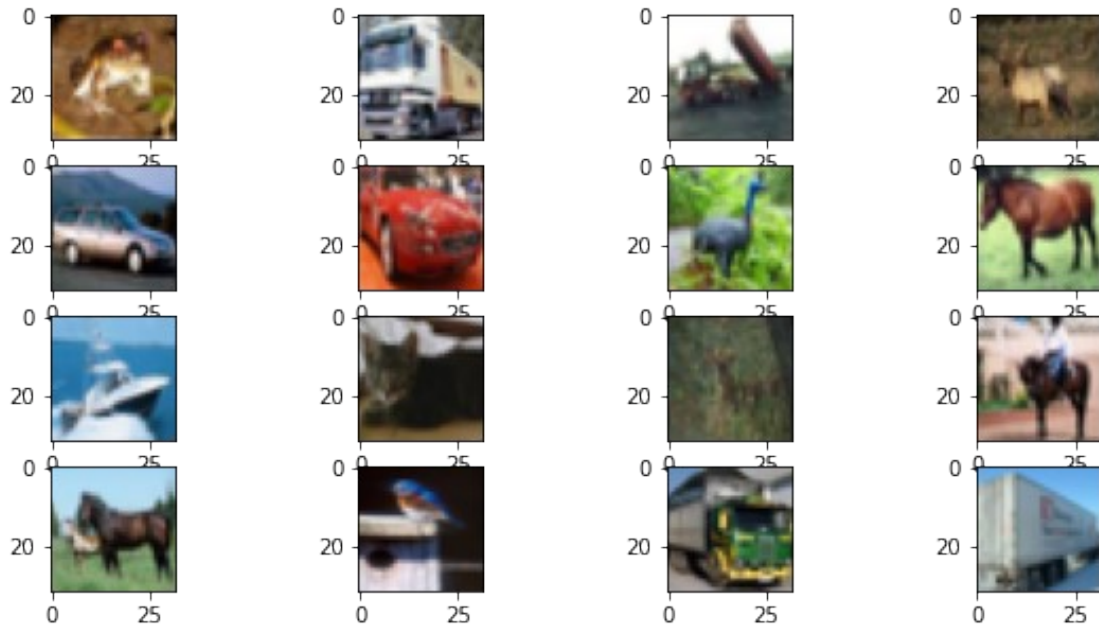
```

Lets See few images

```

fig = plt.figure(figsize=(10, 5))
rows = 4
columns = 4
for i in range(16):
    image = X_train[i]
    fig.add_subplot(rows, columns, i+1)
    plt.imshow(image)
plt.show()

```



Lets Analyse the classes

```
np.unique(y_train)
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

- We have 10 Classes which are already label encoded

```
num_classes = 10
```

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(50000, 32, 32, 3)
```

```
(10000, 32, 32, 3)
```

```
(50000, 1)
```

```
(10000, 1)
```

Lets convert y_train to one hot encoded vector

```
# convert to one hot encoding
```

```
#y_train = tf.keras.utils.to_categorical(y_train, num_classes)
```

```
#y_test = tf.keras.utils.to_categorical(y_test, num_classes)
```

```
# flatten the label values
```

```
y_train, y_test = y_train.flatten(), y_test.flatten()
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

```
(50000,)
```

```
(10000,)
```

We can see that `x_train` values are in range of 0-255 as input is image lets normalise the image so that neural network converge faster

```
X_train = X_train/255.0
```

```
X_test = X_test/255.0
```

using similar architecture provided in reference

```
tf.keras.backend.clear_session()
```

```
# Hyperparameters
```

```
batch_size = 64
```

```
num_classes = 10
```

```
epochs = 50
```

```
l = 40
```

```
num_filter = 12
```

```
compression = 0.5
```

```
dropout_rate = 0.2
```

```
# Dense Block
```

```
def denseblock(input, num_filter = 12, dropout_rate = 0.2):
```

```
    global compression
```

```
    temp = input
```

```
    for _ in range(l):
```

```
        BatchNorm = layers.BatchNormalization()(temp)
```

```
        relu = layers.Activation('relu')(BatchNorm)
```

```
        Conv2D_5_5 = layers.Conv2D(int(num_filter*compression), (5,5),
```

```
        use_bias=True, padding='same')(relu)
```

```
        if dropout_rate>0:
```

```
            Conv2D_5_5 = layers.Dropout(dropout_rate)(Conv2D_5_5)
```

```
        concat = layers.Concatenate(axis=-1)([temp, Conv2D_5_5])
```

```
        temp = concat
```

```
    return temp
```

```
## transition Block
```

```
def transition(input, num_filter = 12, dropout_rate = 0.2):
```

```
    global compression
```

```
    BatchNorm = layers.BatchNormalization()(input)
```

```
    relu = layers.Activation('relu')(BatchNorm)
```

```
    Conv2D_BottleNeck = layers.Conv2D(int(num_filter*compression),  
(5,5), use_bias=False, padding='same')(relu)
```

```
    if dropout_rate>0:
```

```
        Conv2D_BottleNeck = layers.Dropout(dropout_rate)  
(Conv2D_BottleNeck)
```

```
    avg = layers.AveragePooling2D(pool_size=(2,2))(Conv2D_BottleNeck)
```

```
    return avg
```

```
#output layer
```

```
def output_layer(input):
```

```
    global compression
```

```

BatchNorm = layers.BatchNormalization()(input)
relu = layers.Activation('relu')(BatchNorm)
AvgPooling = layers.AveragePooling2D(pool_size=(2,2))(relu)
flat = layers.Flatten()(AvgPooling)
output = layers.Dense(num_classes, activation='softmax')(flat)
return output

```

Lets train the model

```

data_generator = ImageDataGenerator(width_shift_range=0.1,
height_shift_range=0.1, horizontal_flip=True)
# prepare training iterator
train_iterator = data_generator.flow(X_train, y_train,
batch_size=batch_size)

num_filter = 12
dropout_rate = 0
l = 12
input = layers.Input(shape=(img_height, img_width, channel,))
# Image Augmentation
#random_flip = layers.RandomFlip("horizontal")(input)
#random_scaling = layers.Rescaling(scale=1./255)(random_flip)
#random_rotation = layers.RandomRotation(0.4)(random_scaling)
#random_zoom = layers.RandomZoom(0.2,0.2)(random_rotation)
#random_contrast = layers.RandomContrast(0.4)(random_zoom)
#random_brightness = layers.RandomBrightness(0.4,value_range=(0,1))
(random_contrast)
First_Conv2D = layers.Conv2D(32, (3,3),
use_bias=False ,padding='same')(input)

First_Block = denseblock(First_Conv2D, num_filter, dropout_rate)
First_Transition = transition(First_Block, 64, dropout_rate)

Second_Block = denseblock(First_Transition, num_filter, dropout_rate)
Second_Transition = transition(Second_Block, 32, dropout_rate)

Third_Block = denseblock(Second_Transition, num_filter, dropout_rate)
Third_Transition = transition(Third_Block, 32, dropout_rate)

Last_Block = denseblock(Third_Transition, num_filter, dropout_rate)
output = output_layer(Last_Block)

model = Model(inputs=[input], outputs=[output])
model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
Connected to		
=====		

```

=====
input_1 (InputLayer)          [(None, 32, 32, 3)] 0      []

conv2d (Conv2D)               (None, 32, 32, 32) 864
['input_1[0][0]']

batch_normalization (BatchNorm (None, 32, 32, 32) 128
['conv2d[0][0]']
alization)

activation (Activation)        (None, 32, 32, 32) 0
['batch_normalization[0][0]']

conv2d_1 (Conv2D)             (None, 32, 32, 6) 4806
['activation[0][0]']

concatenate (Concatenate)      (None, 32, 32, 38) 0
['conv2d[0][0]',
'conv2d_1[0][0]']

batch_normalization_1 (BatchNo (None, 32, 32, 38) 152
['concatenate[0][0]']
rmalization)

activation_1 (Activation)      (None, 32, 32, 38) 0
['batch_normalization_1[0][0]']

conv2d_2 (Conv2D)             (None, 32, 32, 6) 5706
['activation_1[0][0]']

concatenate_1 (Concatenate)    (None, 32, 32, 44) 0
['concatenate[0][0]',
'conv2d_2[0][0]']

```

batch_normalization_2 (Batch Normalization) ['concatenate_1[0][0]']	(None, 32, 32, 44)	176
activation_2 (Activation) ['batch_normalization_2[0][0]']	(None, 32, 32, 44)	0
conv2d_3 (Conv2D) ['activation_2[0][0]']	(None, 32, 32, 6)	6606
concatenate_2 (Concatenate) ['concatenate_1[0][0]', 'conv2d_3[0][0]']	(None, 32, 32, 50)	0
batch_normalization_3 (Batch Normalization) ['concatenate_2[0][0]']	(None, 32, 32, 50)	200
activation_3 (Activation) ['batch_normalization_3[0][0]']	(None, 32, 32, 50)	0
conv2d_4 (Conv2D) ['activation_3[0][0]']	(None, 32, 32, 6)	7506
concatenate_3 (Concatenate) ['concatenate_2[0][0]', 'conv2d_4[0][0]']	(None, 32, 32, 56)	0
batch_normalization_4 (Batch Normalization) ['concatenate_3[0][0]']	(None, 32, 32, 56)	224
activation_4 (Activation)	(None, 32, 32, 56)	0

['batch_normalization_4[0][0]']

conv2d_5 (Conv2D) (None, 32, 32, 6) 8406
['activation_4[0][0]']

concatenate_4 (Concatenate) (None, 32, 32, 62) 0
['concatenate_3[0][0]',
'conv2d_5[0][0]']

batch_normalization_5 (Batch Normalization) (None, 32, 32, 62) 248
['concatenate_4[0][0]']

activation_5 (Activation) (None, 32, 32, 62) 0
['batch_normalization_5[0][0]']

conv2d_6 (Conv2D) (None, 32, 32, 6) 9306
['activation_5[0][0]']

concatenate_5 (Concatenate) (None, 32, 32, 68) 0
['concatenate_4[0][0]',
'conv2d_6[0][0]']

batch_normalization_6 (Batch Normalization) (None, 32, 32, 68) 272
['concatenate_5[0][0]']

activation_6 (Activation) (None, 32, 32, 68) 0
['batch_normalization_6[0][0]']

conv2d_7 (Conv2D) (None, 32, 32, 6) 10206
['activation_6[0][0]']

concatenate_6 (Concatenate) (None, 32, 32, 74) 0

['concatenate_5[0][0]',

'conv2d_7[0][0]']

batch_normalization_7 (BatchNo (None, 32, 32, 74) 296
['concatenate_6[0][0]']
rmalization)

activation_7 (Activation) (None, 32, 32, 74) 0
['batch_normalization_7[0][0]']

conv2d_8 (Conv2D) (None, 32, 32, 6) 11106
['activation_7[0][0]']

concatenate_7 (Concatenate) (None, 32, 32, 80) 0
['concatenate_6[0][0]',

'conv2d_8[0][0]']

batch_normalization_8 (BatchNo (None, 32, 32, 80) 320
['concatenate_7[0][0]']
rmalization)

activation_8 (Activation) (None, 32, 32, 80) 0
['batch_normalization_8[0][0]']

conv2d_9 (Conv2D) (None, 32, 32, 6) 12006
['activation_8[0][0]']

concatenate_8 (Concatenate) (None, 32, 32, 86) 0
['concatenate_7[0][0]',

'conv2d_9[0][0]']

batch_normalization_9 (BatchNo (None, 32, 32, 86) 344
['concatenate_8[0][0]']
rmalization)

activation_9 (Activation) ['batch_normalization_9[0][0]']	(None, 32, 32, 86)	0
conv2d_10 (Conv2D) ['activation_9[0][0]']	(None, 32, 32, 6)	12906
concatenate_9 (Concatenate) ['concatenate_8[0][0]', 'conv2d_10[0][0]']	(None, 32, 32, 92)	0
batch_normalization_10 (Batch Normalization) ['concatenate_9[0][0]']	(None, 32, 32, 92)	368
activation_10 (Activation) ['batch_normalization_10[0][0]']	(None, 32, 32, 92)	0
conv2d_11 (Conv2D) ['activation_10[0][0]']	(None, 32, 32, 6)	13806
concatenate_10 (Concatenate) ['concatenate_9[0][0]', 'conv2d_11[0][0]']	(None, 32, 32, 98)	0
batch_normalization_11 (Batch Normalization) ['concatenate_10[0][0]']	(None, 32, 32, 98)	392
activation_11 (Activation) ['batch_normalization_11[0][0]']	(None, 32, 32, 98)	0
conv2d_12 (Conv2D)	(None, 32, 32, 6)	14706

['activation_11[0][0]']

concatenate_11 (Concatenate) (None, 32, 32, 104) 0
['concatenate_10[0][0]',

'conv2d_12[0][0]']

batch_normalization_12 (BatchN (None, 32, 32, 104) 416
['concatenate_11[0][0]']
ormalization)

activation_12 (Activation) (None, 32, 32, 104) 0
['batch_normalization_12[0][0]']

conv2d_13 (Conv2D) (None, 32, 32, 32) 83200
['activation_12[0][0]']

average_pooling2d (AveragePool (None, 16, 16, 32) 0
['conv2d_13[0][0]']
ing2D)

batch_normalization_13 (BatchN (None, 16, 16, 32) 128
['average_pooling2d[0][0]']
ormalization)

activation_13 (Activation) (None, 16, 16, 32) 0
['batch_normalization_13[0][0]']

conv2d_14 (Conv2D) (None, 16, 16, 6) 4806
['activation_13[0][0]']

concatenate_12 (Concatenate) (None, 16, 16, 38) 0
['average_pooling2d[0][0]',

'conv2d_14[0][0]']

batch_normalization_14 (Batch Normalization)	(None, 16, 16, 38)	152
['concatenate_12[0][0]']		
activation_14 (Activation)	(None, 16, 16, 38)	0
['batch_normalization_14[0][0]']		
conv2d_15 (Conv2D)	(None, 16, 16, 6)	5706
['activation_14[0][0]']		
concatenate_13 (Concatenate)	(None, 16, 16, 44)	0
['concatenate_12[0][0]',		
'conv2d_15[0][0]']		
batch_normalization_15 (Batch Normalization)	(None, 16, 16, 44)	176
['concatenate_13[0][0]']		
activation_15 (Activation)	(None, 16, 16, 44)	0
['batch_normalization_15[0][0]']		
conv2d_16 (Conv2D)	(None, 16, 16, 6)	6606
['activation_15[0][0]']		
concatenate_14 (Concatenate)	(None, 16, 16, 50)	0
['concatenate_13[0][0]',		
'conv2d_16[0][0]']		
batch_normalization_16 (Batch Normalization)	(None, 16, 16, 50)	200
['concatenate_14[0][0]']		
activation_16 (Activation)	(None, 16, 16, 50)	0

['batch_normalization_16[0][0]']

conv2d_17 (Conv2D) (None, 16, 16, 6) 7506
['activation_16[0][0]']

concatenate_15 (Concatenate) (None, 16, 16, 56) 0
['concatenate_14[0][0]',
'conv2d_17[0][0]']

batch_normalization_17 (BatchN (None, 16, 16, 56) 224
['concatenate_15[0][0]']
ormalization)

activation_17 (Activation) (None, 16, 16, 56) 0
['batch_normalization_17[0][0]']

conv2d_18 (Conv2D) (None, 16, 16, 6) 8406
['activation_17[0][0]']

concatenate_16 (Concatenate) (None, 16, 16, 62) 0
['concatenate_15[0][0]',
'conv2d_18[0][0]']

batch_normalization_18 (BatchN (None, 16, 16, 62) 248
['concatenate_16[0][0]']
ormalization)

activation_18 (Activation) (None, 16, 16, 62) 0
['batch_normalization_18[0][0]']

conv2d_19 (Conv2D) (None, 16, 16, 6) 9306
['activation_18[0][0]']

concatenate_17 (Concatenate) (None, 16, 16, 68) 0

['concatenate_16[0][0]',

'conv2d_19[0][0]']

batch_normalization_19 (BatchN (None, 16, 16, 68) 272
['concatenate_17[0][0]']
ormalization)

activation_19 (Activation) (None, 16, 16, 68) 0
['batch_normalization_19[0][0]']

conv2d_20 (Conv2D) (None, 16, 16, 6) 10206
['activation_19[0][0]']

concatenate_18 (Concatenate) (None, 16, 16, 74) 0
['concatenate_17[0][0]',

'conv2d_20[0][0]']

batch_normalization_20 (BatchN (None, 16, 16, 74) 296
['concatenate_18[0][0]']
ormalization)

activation_20 (Activation) (None, 16, 16, 74) 0
['batch_normalization_20[0][0]']

conv2d_21 (Conv2D) (None, 16, 16, 6) 11106
['activation_20[0][0]']

concatenate_19 (Concatenate) (None, 16, 16, 80) 0
['concatenate_18[0][0]',

'conv2d_21[0][0]']

batch_normalization_21 (BatchN (None, 16, 16, 80) 320
['concatenate_19[0][0]']
ormalization)

activation_21 (Activation)	(None, 16, 16, 80)	0
['batch_normalization_21[0][0]']		
conv2d_22 (Conv2D)	(None, 16, 16, 6)	12006
['activation_21[0][0]']		
concatenate_20 (Concatenate)	(None, 16, 16, 86)	0
['concatenate_19[0][0]',		
'conv2d_22[0][0]']		
batch_normalization_22 (BatchN	(None, 16, 16, 86)	344
['concatenate_20[0][0]']		
ormalization)		
activation_22 (Activation)	(None, 16, 16, 86)	0
['batch_normalization_22[0][0]']		
conv2d_23 (Conv2D)	(None, 16, 16, 6)	12906
['activation_22[0][0]']		
concatenate_21 (Concatenate)	(None, 16, 16, 92)	0
['concatenate_20[0][0]',		
'conv2d_23[0][0]']		
batch_normalization_23 (BatchN	(None, 16, 16, 92)	368
['concatenate_21[0][0]']		
ormalization)		
activation_23 (Activation)	(None, 16, 16, 92)	0
['batch_normalization_23[0][0]']		
conv2d_24 (Conv2D)	(None, 16, 16, 6)	13806

['activation_23[0][0]']

concatenate_22 (Concatenate) (None, 16, 16, 98) 0
['concatenate_21[0][0]',

'conv2d_24[0][0]']

batch_normalization_24 (BatchN (None, 16, 16, 98) 392
['concatenate_22[0][0]']
ormalization)

activation_24 (Activation) (None, 16, 16, 98) 0
['batch_normalization_24[0][0]']

conv2d_25 (Conv2D) (None, 16, 16, 6) 14706
['activation_24[0][0]']

concatenate_23 (Concatenate) (None, 16, 16, 104) 0
['concatenate_22[0][0]',

'conv2d_25[0][0]']

batch_normalization_25 (BatchN (None, 16, 16, 104) 416
['concatenate_23[0][0]']
ormalization)

activation_25 (Activation) (None, 16, 16, 104) 0
['batch_normalization_25[0][0]']

conv2d_26 (Conv2D) (None, 16, 16, 16) 41600
['activation_25[0][0]']

average_pooling2d_1 (AveragePo (None, 8, 8, 16) 0
['conv2d_26[0][0]']
oling2D)

batch_normalization_26 (Batch Normalization) ['average_pooling2d_1[0][0]']	(None, 8, 8, 16)	64
activation_26 (Activation) ['batch_normalization_26[0][0]']	(None, 8, 8, 16)	0
conv2d_27 (Conv2D) ['activation_26[0][0]']	(None, 8, 8, 6)	2406
concatenate_24 (Concatenate) ['average_pooling2d_1[0][0]', 'conv2d_27[0][0]']	(None, 8, 8, 22)	0
batch_normalization_27 (Batch Normalization) ['concatenate_24[0][0]']	(None, 8, 8, 22)	88
activation_27 (Activation) ['batch_normalization_27[0][0]']	(None, 8, 8, 22)	0
conv2d_28 (Conv2D) ['activation_27[0][0]']	(None, 8, 8, 6)	3306
concatenate_25 (Concatenate) ['concatenate_24[0][0]', 'conv2d_28[0][0]']	(None, 8, 8, 28)	0
batch_normalization_28 (Batch Normalization) ['concatenate_25[0][0]']	(None, 8, 8, 28)	112
activation_28 (Activation)	(None, 8, 8, 28)	0

['batch_normalization_28[0][0]']

conv2d_29 (Conv2D)	(None, 8, 8, 6)	4206
['activation_28[0][0]']		

concatenate_26 (Concatenate)	(None, 8, 8, 34)	0
['concatenate_25[0][0]',		
'conv2d_29[0][0]']		

batch_normalization_29 (BatchN	(None, 8, 8, 34)	136
['concatenate_26[0][0]']		
ormalization)		

activation_29 (Activation)	(None, 8, 8, 34)	0
['batch_normalization_29[0][0]']		

conv2d_30 (Conv2D)	(None, 8, 8, 6)	5106
['activation_29[0][0]']		

concatenate_27 (Concatenate)	(None, 8, 8, 40)	0
['concatenate_26[0][0]',		
'conv2d_30[0][0]']		

batch_normalization_30 (BatchN	(None, 8, 8, 40)	160
['concatenate_27[0][0]']		
ormalization)		

activation_30 (Activation)	(None, 8, 8, 40)	0
['batch_normalization_30[0][0]']		

conv2d_31 (Conv2D)	(None, 8, 8, 6)	6006
['activation_30[0][0]']		

concatenate_28 (Concatenate)	(None, 8, 8, 46)	0
------------------------------	------------------	---

['concatenate_27[0][0]',

'conv2d_31[0][0]']

batch_normalization_31 (BatchN (None, 8, 8, 46) 184
['concatenate_28[0][0]']
ormalization)

activation_31 (Activation) (None, 8, 8, 46) 0
['batch_normalization_31[0][0]']

conv2d_32 (Conv2D) (None, 8, 8, 6) 6906
['activation_31[0][0]']

concatenate_29 (Concatenate) (None, 8, 8, 52) 0
['concatenate_28[0][0]',
'conv2d_32[0][0]']

batch_normalization_32 (BatchN (None, 8, 8, 52) 208
['concatenate_29[0][0]']
ormalization)

activation_32 (Activation) (None, 8, 8, 52) 0
['batch_normalization_32[0][0]']

conv2d_33 (Conv2D) (None, 8, 8, 6) 7806
['activation_32[0][0]']

concatenate_30 (Concatenate) (None, 8, 8, 58) 0
['concatenate_29[0][0]',
'conv2d_33[0][0]']

batch_normalization_33 (BatchN (None, 8, 8, 58) 232
['concatenate_30[0][0]']
ormalization)

activation_33 (Activation) ['batch_normalization_33[0][0]']	(None, 8, 8, 58)	0
conv2d_34 (Conv2D) ['activation_33[0][0]']	(None, 8, 8, 6)	8706
concatenate_31 (Concatenate) ['concatenate_30[0][0]', 'conv2d_34[0][0]']	(None, 8, 8, 64)	0
batch_normalization_34 (BatchN ['concatenate_31[0][0]'] ormalization)	(None, 8, 8, 64)	256
activation_34 (Activation) ['batch_normalization_34[0][0]']	(None, 8, 8, 64)	0
conv2d_35 (Conv2D) ['activation_34[0][0]']	(None, 8, 8, 6)	9606
concatenate_32 (Concatenate) ['concatenate_31[0][0]', 'conv2d_35[0][0]']	(None, 8, 8, 70)	0
batch_normalization_35 (BatchN ['concatenate_32[0][0]'] ormalization)	(None, 8, 8, 70)	280
activation_35 (Activation) ['batch_normalization_35[0][0]']	(None, 8, 8, 70)	0
conv2d_36 (Conv2D)	(None, 8, 8, 6)	10506

['activation_35[0][0]']

concatenate_33 (Concatenate) (None, 8, 8, 76) 0
['concatenate_32[0][0]',

'conv2d_36[0][0]']

batch_normalization_36 (BatchN (None, 8, 8, 76) 304
['concatenate_33[0][0]']
ormalization)

activation_36 (Activation) (None, 8, 8, 76) 0
['batch_normalization_36[0][0]']

conv2d_37 (Conv2D) (None, 8, 8, 6) 11406
['activation_36[0][0]']

concatenate_34 (Concatenate) (None, 8, 8, 82) 0
['concatenate_33[0][0]',

'conv2d_37[0][0]']

batch_normalization_37 (BatchN (None, 8, 8, 82) 328
['concatenate_34[0][0]']
ormalization)

activation_37 (Activation) (None, 8, 8, 82) 0
['batch_normalization_37[0][0]']

conv2d_38 (Conv2D) (None, 8, 8, 6) 12306
['activation_37[0][0]']

concatenate_35 (Concatenate) (None, 8, 8, 88) 0
['concatenate_34[0][0]',

'conv2d_38[0][0]']

batch_normalization_38 (Batch Normalization) ['concatenate_35[0][0]']	(None, 8, 8, 88)	352
activation_38 (Activation) ['batch_normalization_38[0][0]']	(None, 8, 8, 88)	0
conv2d_39 (Conv2D) ['activation_38[0][0]']	(None, 8, 8, 16)	35200
average_pooling2d_2 (Average Pooling2D) ['conv2d_39[0][0]']	(None, 4, 4, 16)	0
batch_normalization_39 (Batch Normalization) ['average_pooling2d_2[0][0]']	(None, 4, 4, 16)	64
activation_39 (Activation) ['batch_normalization_39[0][0]']	(None, 4, 4, 16)	0
conv2d_40 (Conv2D) ['activation_39[0][0]']	(None, 4, 4, 6)	2406
concatenate_36 (Concatenate) ['average_pooling2d_2[0][0]', 'conv2d_40[0][0]']	(None, 4, 4, 22)	0
batch_normalization_40 (Batch Normalization) ['concatenate_36[0][0]']	(None, 4, 4, 22)	88
activation_40 (Activation)	(None, 4, 4, 22)	0

['batch_normalization_40[0][0]']

conv2d_41 (Conv2D) (None, 4, 4, 6) 3306
['activation_40[0][0]']

concatenate_37 (Concatenate) (None, 4, 4, 28) 0
['concatenate_36[0][0]',
'conv2d_41[0][0]']

batch_normalization_41 (BatchN (None, 4, 4, 28) 112
['concatenate_37[0][0]']
ormalization)

activation_41 (Activation) (None, 4, 4, 28) 0
['batch_normalization_41[0][0]']

conv2d_42 (Conv2D) (None, 4, 4, 6) 4206
['activation_41[0][0]']

concatenate_38 (Concatenate) (None, 4, 4, 34) 0
['concatenate_37[0][0]',
'conv2d_42[0][0]']

batch_normalization_42 (BatchN (None, 4, 4, 34) 136
['concatenate_38[0][0]']
ormalization)

activation_42 (Activation) (None, 4, 4, 34) 0
['batch_normalization_42[0][0]']

conv2d_43 (Conv2D) (None, 4, 4, 6) 5106
['activation_42[0][0]']

concatenate_39 (Concatenate) (None, 4, 4, 40) 0

['concatenate_38[0][0]',

'conv2d_43[0][0]']

batch_normalization_43 (BatchN (None, 4, 4, 40) 160
['concatenate_39[0][0]']
ormalization)

activation_43 (Activation) (None, 4, 4, 40) 0
['batch_normalization_43[0][0]']

conv2d_44 (Conv2D) (None, 4, 4, 6) 6006
['activation_43[0][0]']

concatenate_40 (Concatenate) (None, 4, 4, 46) 0
['concatenate_39[0][0]',

'conv2d_44[0][0]']

batch_normalization_44 (BatchN (None, 4, 4, 46) 184
['concatenate_40[0][0]']
ormalization)

activation_44 (Activation) (None, 4, 4, 46) 0
['batch_normalization_44[0][0]']

conv2d_45 (Conv2D) (None, 4, 4, 6) 6906
['activation_44[0][0]']

concatenate_41 (Concatenate) (None, 4, 4, 52) 0
['concatenate_40[0][0]',

'conv2d_45[0][0]']

batch_normalization_45 (BatchN (None, 4, 4, 52) 208
['concatenate_41[0][0]']
ormalization)

activation_45 (Activation) ['batch_normalization_45[0][0]']	(None, 4, 4, 52)	0
conv2d_46 (Conv2D) ['activation_45[0][0]']	(None, 4, 4, 6)	7806
concatenate_42 (Concatenate) ['concatenate_41[0][0]', 'conv2d_46[0][0]']	(None, 4, 4, 58)	0
batch_normalization_46 (BatchN ['concatenate_42[0][0]'] ormalization)	(None, 4, 4, 58)	232
activation_46 (Activation) ['batch_normalization_46[0][0]']	(None, 4, 4, 58)	0
conv2d_47 (Conv2D) ['activation_46[0][0]']	(None, 4, 4, 6)	8706
concatenate_43 (Concatenate) ['concatenate_42[0][0]', 'conv2d_47[0][0]']	(None, 4, 4, 64)	0
batch_normalization_47 (BatchN ['concatenate_43[0][0]'] ormalization)	(None, 4, 4, 64)	256
activation_47 (Activation) ['batch_normalization_47[0][0]']	(None, 4, 4, 64)	0
conv2d_48 (Conv2D)	(None, 4, 4, 6)	9606

['activation_47[0][0]']

concatenate_44 (Concatenate) (None, 4, 4, 70) 0
['concatenate_43[0][0]',

'conv2d_48[0][0]']

batch_normalization_48 (BatchN (None, 4, 4, 70) 280
['concatenate_44[0][0]']
ormalization)

activation_48 (Activation) (None, 4, 4, 70) 0
['batch_normalization_48[0][0]']

conv2d_49 (Conv2D) (None, 4, 4, 6) 10506
['activation_48[0][0]']

concatenate_45 (Concatenate) (None, 4, 4, 76) 0
['concatenate_44[0][0]',

'conv2d_49[0][0]']

batch_normalization_49 (BatchN (None, 4, 4, 76) 304
['concatenate_45[0][0]']
ormalization)

activation_49 (Activation) (None, 4, 4, 76) 0
['batch_normalization_49[0][0]']

conv2d_50 (Conv2D) (None, 4, 4, 6) 11406
['activation_49[0][0]']

concatenate_46 (Concatenate) (None, 4, 4, 82) 0
['concatenate_45[0][0]',

'conv2d_50[0][0]']

batch_normalization_50 (Batch Normalization) ['concatenate_46[0][0]']	(None, 4, 4, 82)	328
activation_50 (Activation) ['batch_normalization_50[0][0]']	(None, 4, 4, 82)	0
conv2d_51 (Conv2D) ['activation_50[0][0]']	(None, 4, 4, 6)	12306
concatenate_47 (Concatenate) ['concatenate_46[0][0]', 'conv2d_51[0][0]']	(None, 4, 4, 88)	0
batch_normalization_51 (Batch Normalization) ['concatenate_47[0][0]']	(None, 4, 4, 88)	352
activation_51 (Activation) ['batch_normalization_51[0][0]']	(None, 4, 4, 88)	0
average_pooling2d_3 (Average Pooling2D) ['activation_51[0][0]']	(None, 2, 2, 88)	0
flatten (Flatten) ['average_pooling2d_3[0][0]']	(None, 352)	0
dense (Dense) ['flatten[0][0]']	(None, 10)	3530

=====

=====

Total params: 587,562

Trainable params: 581,322
Non-trainable params: 6,240

```
print(len(model.layers))
```

211

```
filepath="model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"  
checkpoint = ModelCheckpoint(filepath=filepath,  
monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')
```

```
# Load the TensorBoard notebook extension
```

```
%load_ext tensorboard
```

```
log_dir = os.path.join("logs", 'fits',  
datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
```

```
tensorboard_callback =  
tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_  
graph=True)
```

```
%reload_ext tensorboard
```

```
# determine Loss function and Optimizer
```

```
model.compile(loss='sparse_categorical_crossentropy',  
optimizer=Adam(), metrics=['accuracy'])
```

```
reduce_lr = ReduceLROnPlateau(monitor='val_accuracy',  
factor=0.1, patience=5, min_lr=0.000001)
```

```
model.fit(train_iterator, batch_size=batch_size, epochs=5, verbose=1, vali  
dation_data=(X_test, y_test), callbacks =  
[checkpoint, reduce_lr, tensorboard_callback])
```

Epoch 1/5

```
782/782 [=====] - ETA: 0s - loss: 0.2645 -  
accuracy: 0.9068
```

Epoch 1: val_accuracy did not improve from 0.87340

```
782/782 [=====] - 84s 108ms/step - loss:  
0.2645 - accuracy: 0.9068 - val_loss: 0.4605 - val_accuracy: 0.8552 -  
lr: 0.0010
```

Epoch 2/5

```
782/782 [=====] - ETA: 0s - loss: 0.2440 -  
accuracy: 0.9146
```

Epoch 2: val_accuracy did not improve from 0.87340

```
782/782 [=====] - 85s 108ms/step - loss:  
0.2440 - accuracy: 0.9146 - val_loss: 0.6535 - val_accuracy: 0.8000 -  
lr: 0.0010
```

Epoch 3/5

```
782/782 [=====] - ETA: 0s - loss: 0.2338 -  
accuracy: 0.9174
```

Epoch 3: val_accuracy did not improve from 0.87340

```
782/782 [=====] - 84s 107ms/step - loss:
```

0.2338 - accuracy: 0.9174 - val_loss: 0.4759 - val_accuracy: 0.8495 -
lr: 0.0010

Epoch 4/5

782/782 [=====] - ETA: 0s - loss: 0.2297 -
accuracy: 0.9197

Epoch 4: val_accuracy did not improve from 0.87340

782/782 [=====] - 84s 107ms/step - loss:
0.2297 - accuracy: 0.9197 - val_loss: 0.6275 - val_accuracy: 0.8142 -
lr: 0.0010

Epoch 5/5

782/782 [=====] - ETA: 0s - loss: 0.2190 -
accuracy: 0.9226

Epoch 5: val_accuracy did not improve from 0.87340

782/782 [=====] - 84s 107ms/step - loss:
0.2190 - accuracy: 0.9226 - val_loss: 0.4594 - val_accuracy: 0.8588 -
lr: 0.0010

<keras.callbacks.History at 0x7fef046bd00>

```
model.fit(train_iterator, batch_size=batch_size, epochs=15, verbose=1, val  
          idation_data=(X_test, y_test), callbacks =  
          [checkpoint, reduce_lr, tensorboard_callback])
```

Epoch 1/15

782/782 [=====] - ETA: 0s - loss: 0.2113 -
accuracy: 0.9255

Epoch 1: val_accuracy did not improve from 0.87340

782/782 [=====] - 84s 107ms/step - loss:
0.2113 - accuracy: 0.9255 - val_loss: 0.5140 - val_accuracy: 0.8424 -
lr: 0.0010

Epoch 2/15

782/782 [=====] - ETA: 0s - loss: 0.2101 -
accuracy: 0.9261

Epoch 2: val_accuracy did not improve from 0.87340

782/782 [=====] - 84s 108ms/step - loss:
0.2101 - accuracy: 0.9261 - val_loss: 0.5460 - val_accuracy: 0.8392 -
lr: 0.0010

Epoch 3/15

782/782 [=====] - ETA: 0s - loss: 0.2022 -
accuracy: 0.9288

Epoch 3: val_accuracy did not improve from 0.87340

782/782 [=====] - 85s 109ms/step - loss:
0.2022 - accuracy: 0.9288 - val_loss: 0.4144 - val_accuracy: 0.8724 -
lr: 0.0010

Epoch 4/15

782/782 [=====] - ETA: 0s - loss: 0.1981 -
accuracy: 0.9305

Epoch 4: val_accuracy did not improve from 0.87340

782/782 [=====] - 84s 108ms/step - loss:
0.1981 - accuracy: 0.9305 - val_loss: 0.5668 - val_accuracy: 0.8364 -
lr: 0.0010

Epoch 5/15
782/782 [=====] - ETA: 0s - loss: 0.1964 - accuracy: 0.9294
Epoch 5: val_accuracy did not improve from 0.87340
782/782 [=====] - 84s 107ms/step - loss: 0.1964 - accuracy: 0.9294 - val_loss: 0.4441 - val_accuracy: 0.8627 - lr: 0.0010
Epoch 6/15
782/782 [=====] - ETA: 0s - loss: 0.1842 - accuracy: 0.9354
Epoch 6: val_accuracy did not improve from 0.87340
782/782 [=====] - 84s 107ms/step - loss: 0.1842 - accuracy: 0.9354 - val_loss: 0.5729 - val_accuracy: 0.8386 - lr: 0.0010
Epoch 7/15
782/782 [=====] - ETA: 0s - loss: 0.1840 - accuracy: 0.9349
Epoch 7: val_accuracy did not improve from 0.87340
782/782 [=====] - 84s 108ms/step - loss: 0.1840 - accuracy: 0.9349 - val_loss: 0.4727 - val_accuracy: 0.8580 - lr: 0.0010
Epoch 8/15
782/782 [=====] - ETA: 0s - loss: 0.1779 - accuracy: 0.9361
Epoch 8: val_accuracy did not improve from 0.87340
782/782 [=====] - 83s 106ms/step - loss: 0.1779 - accuracy: 0.9361 - val_loss: 0.5875 - val_accuracy: 0.8392 - lr: 0.0010
Epoch 9/15
782/782 [=====] - ETA: 0s - loss: 0.1319 - accuracy: 0.9549
Epoch 9: val_accuracy improved from 0.87340 to 0.89210, saving model to model_save/weights-09-0.8921.hdf5
782/782 [=====] - 84s 107ms/step - loss: 0.1319 - accuracy: 0.9549 - val_loss: 0.3449 - val_accuracy: 0.8921 - lr: 1.0000e-04
Epoch 10/15
782/782 [=====] - ETA: 0s - loss: 0.1096 - accuracy: 0.9628
Epoch 10: val_accuracy improved from 0.89210 to 0.89370, saving model to model_save/weights-10-0.8937.hdf5
782/782 [=====] - 85s 109ms/step - loss: 0.1096 - accuracy: 0.9628 - val_loss: 0.3468 - val_accuracy: 0.8937 - lr: 1.0000e-04
Epoch 11/15
782/782 [=====] - ETA: 0s - loss: 0.1026 - accuracy: 0.9641
Epoch 11: val_accuracy improved from 0.89370 to 0.89480, saving model to model_save/weights-11-0.8948.hdf5
782/782 [=====] - 85s 108ms/step - loss:

0.1026 - accuracy: 0.9641 - val_loss: 0.3530 - val_accuracy: 0.8948 -
lr: 1.0000e-04

Epoch 12/15

782/782 [=====] - ETA: 0s - loss: 0.0959 -
accuracy: 0.9669

Epoch 12: val_accuracy improved from 0.89480 to 0.89490, saving model
to model_save/weights-12-0.8949.hdf5

782/782 [=====] - 84s 107ms/step - loss:
0.0959 - accuracy: 0.9669 - val_loss: 0.3550 - val_accuracy: 0.8949 -
lr: 1.0000e-04

Epoch 13/15

782/782 [=====] - ETA: 0s - loss: 0.0928 -
accuracy: 0.9680

Epoch 13: val_accuracy improved from 0.89490 to 0.89510, saving model
to model_save/weights-13-0.8951.hdf5

782/782 [=====] - 85s 108ms/step - loss:
0.0928 - accuracy: 0.9680 - val_loss: 0.3664 - val_accuracy: 0.8951 -
lr: 1.0000e-04

Epoch 14/15

782/782 [=====] - ETA: 0s - loss: 0.0888 -
accuracy: 0.9696

Epoch 14: val_accuracy did not improve from 0.89510

782/782 [=====] - 83s 106ms/step - loss:
0.0888 - accuracy: 0.9696 - val_loss: 0.3708 - val_accuracy: 0.8927 -
lr: 1.0000e-04

Epoch 15/15

782/782 [=====] - ETA: 0s - loss: 0.0891 -
accuracy: 0.9691

Epoch 15: val_accuracy improved from 0.89510 to 0.89610, saving model
to model_save/weights-15-0.8961.hdf5

782/782 [=====] - 86s 109ms/step - loss:
0.0891 - accuracy: 0.9691 - val_loss: 0.3680 - val_accuracy: 0.8961 -
lr: 1.0000e-04

<keras.callbacks.History at 0x7fef5de97970>

```
model.fit(train_iterator, batch_size=batch_size, epochs=15, verbose=1, val  
idation_data=(X_test, y_test), callbacks =  
[checkpoint, reduce_lr, tensorboard_callback])
```

Epoch 1/15

782/782 [=====] - ETA: 0s - loss: 0.0844 -
accuracy: 0.9709

Epoch 1: val_accuracy did not improve from 0.89610

782/782 [=====] - 85s 108ms/step - loss:
0.0844 - accuracy: 0.9709 - val_loss: 0.3761 - val_accuracy: 0.8957 -
lr: 1.0000e-04

Epoch 2/15

782/782 [=====] - ETA: 0s - loss: 0.0838 -
accuracy: 0.9709

Epoch 2: val_accuracy did not improve from 0.89610

782/782 [=====] - 83s 106ms/step - loss: 0.0838 - accuracy: 0.9709 - val_loss: 0.3831 - val_accuracy: 0.8939 - lr: 1.0000e-04
Epoch 3/15
782/782 [=====] - ETA: 0s - loss: 0.0803 - accuracy: 0.9721
Epoch 3: val_accuracy did not improve from 0.89610
782/782 [=====] - 84s 107ms/step - loss: 0.0803 - accuracy: 0.9721 - val_loss: 0.3754 - val_accuracy: 0.8952 - lr: 1.0000e-04
Epoch 4/15
782/782 [=====] - ETA: 0s - loss: 0.0796 - accuracy: 0.9720
Epoch 4: val_accuracy did not improve from 0.89610
782/782 [=====] - 83s 106ms/step - loss: 0.0796 - accuracy: 0.9720 - val_loss: 0.3859 - val_accuracy: 0.8921 - lr: 1.0000e-04
Epoch 5/15
782/782 [=====] - ETA: 0s - loss: 0.0789 - accuracy: 0.9721
Epoch 5: val_accuracy improved from 0.89610 to 0.89650, saving model to model_save/weights-05-0.8965.hdf5
782/782 [=====] - 85s 109ms/step - loss: 0.0789 - accuracy: 0.9721 - val_loss: 0.3807 - val_accuracy: 0.8965 - lr: 1.0000e-04
Epoch 6/15
782/782 [=====] - ETA: 0s - loss: 0.0763 - accuracy: 0.9732
Epoch 6: val_accuracy did not improve from 0.89650
782/782 [=====] - 84s 107ms/step - loss: 0.0763 - accuracy: 0.9732 - val_loss: 0.3893 - val_accuracy: 0.8940 - lr: 1.0000e-04
Epoch 7/15
782/782 [=====] - ETA: 0s - loss: 0.0725 - accuracy: 0.9747
Epoch 7: val_accuracy did not improve from 0.89650
782/782 [=====] - 83s 106ms/step - loss: 0.0725 - accuracy: 0.9747 - val_loss: 0.4120 - val_accuracy: 0.8926 - lr: 1.0000e-04
Epoch 8/15
782/782 [=====] - ETA: 0s - loss: 0.0715 - accuracy: 0.9754
Epoch 8: val_accuracy did not improve from 0.89650
782/782 [=====] - 84s 107ms/step - loss: 0.0715 - accuracy: 0.9754 - val_loss: 0.3988 - val_accuracy: 0.8942 - lr: 1.0000e-04
Epoch 9/15
782/782 [=====] - ETA: 0s - loss: 0.0687 - accuracy: 0.9755
Epoch 9: val_accuracy did not improve from 0.89650

782/782 [=====] - 83s 107ms/step - loss: 0.0687 - accuracy: 0.9755 - val_loss: 0.4002 - val_accuracy: 0.8959 - lr: 1.0000e-04
Epoch 10/15
782/782 [=====] - ETA: 0s - loss: 0.0691 - accuracy: 0.9759
Epoch 10: val_accuracy did not improve from 0.89650
782/782 [=====] - 83s 107ms/step - loss: 0.0691 - accuracy: 0.9759 - val_loss: 0.4074 - val_accuracy: 0.8932 - lr: 1.0000e-04
Epoch 11/15
782/782 [=====] - ETA: 0s - loss: 0.0660 - accuracy: 0.9766
Epoch 11: val_accuracy did not improve from 0.89650
782/782 [=====] - 86s 111ms/step - loss: 0.0660 - accuracy: 0.9766 - val_loss: 0.3998 - val_accuracy: 0.8952 - lr: 1.0000e-05
Epoch 12/15
782/782 [=====] - ETA: 0s - loss: 0.0638 - accuracy: 0.9774
Epoch 12: val_accuracy did not improve from 0.89650
782/782 [=====] - 85s 108ms/step - loss: 0.0638 - accuracy: 0.9774 - val_loss: 0.3996 - val_accuracy: 0.8962 - lr: 1.0000e-05
Epoch 13/15
782/782 [=====] - ETA: 0s - loss: 0.0631 - accuracy: 0.9787
Epoch 13: val_accuracy did not improve from 0.89650
782/782 [=====] - 83s 106ms/step - loss: 0.0631 - accuracy: 0.9787 - val_loss: 0.3979 - val_accuracy: 0.8964 - lr: 1.0000e-05
Epoch 14/15
782/782 [=====] - ETA: 0s - loss: 0.0637 - accuracy: 0.9779
Epoch 14: val_accuracy improved from 0.89650 to 0.89670, saving model to model_save/weights-14-0.8967.hdf5
782/782 [=====] - 84s 108ms/step - loss: 0.0637 - accuracy: 0.9779 - val_loss: 0.3978 - val_accuracy: 0.8967 - lr: 1.0000e-05
Epoch 15/15
782/782 [=====] - ETA: 0s - loss: 0.0627 - accuracy: 0.9784
Epoch 15: val_accuracy improved from 0.89670 to 0.89750, saving model to model_save/weights-15-0.8975.hdf5
782/782 [=====] - 85s 108ms/step - loss: 0.0627 - accuracy: 0.9784 - val_loss: 0.3961 - val_accuracy: 0.8975 - lr: 1.0000e-05

<keras.callbacks.History at 0x7fef5deb02b0>

```
data_generator = ImageDataGenerator(width_shift_range=0.1,
height_shift_range=0.1, rotation_range =
0.2,shear_range=0.1,zoom_range=0.1,horizontal_flip=True)
# prepare training iterator
train_iterator = data_generator.flow(X_train, y_train,
batch_size=batch_size)

model.fit(train_iterator,batch_size=batch_size,epochs=10,verbose=1,validation_data=(X_test, y_test),callbacks =
[checkpoint,reduce_lr,tensorboard_callback])
```

Epoch 1/10

782/782 [=====] - ETA: 0s - loss: 0.0865 - accuracy: 0.9695

Epoch 1: val_accuracy did not improve from 0.89750

782/782 [=====] - 85s 109ms/step - loss: 0.0865 - accuracy: 0.9695 - val_loss: 0.4057 - val_accuracy: 0.8952 - lr: 1.0000e-05

Epoch 2/10

782/782 [=====] - ETA: 0s - loss: 0.0836 - accuracy: 0.9704

Epoch 2: val_accuracy did not improve from 0.89750

782/782 [=====] - 84s 108ms/step - loss: 0.0836 - accuracy: 0.9704 - val_loss: 0.4039 - val_accuracy: 0.8956 - lr: 1.0000e-05

Epoch 3/10

782/782 [=====] - ETA: 0s - loss: 0.0841 - accuracy: 0.9695

Epoch 3: val_accuracy did not improve from 0.89750

782/782 [=====] - 84s 108ms/step - loss: 0.0841 - accuracy: 0.9695 - val_loss: 0.4051 - val_accuracy: 0.8946 - lr: 1.0000e-05

Epoch 4/10

782/782 [=====] - ETA: 0s - loss: 0.0813 - accuracy: 0.9710

Epoch 4: val_accuracy did not improve from 0.89750

782/782 [=====] - 84s 108ms/step - loss: 0.0813 - accuracy: 0.9710 - val_loss: 0.4051 - val_accuracy: 0.8951 - lr: 1.0000e-05

Epoch 5/10

782/782 [=====] - ETA: 0s - loss: 0.0835 - accuracy: 0.9701

Epoch 5: val_accuracy did not improve from 0.89750

782/782 [=====] - 84s 107ms/step - loss: 0.0835 - accuracy: 0.9701 - val_loss: 0.4027 - val_accuracy: 0.8946 - lr: 1.0000e-05

Epoch 6/10

782/782 [=====] - ETA: 0s - loss: 0.0818 - accuracy: 0.9714

Epoch 6: val_accuracy did not improve from 0.89750

782/782 [=====] - 85s 109ms/step - loss:

0.0818 - accuracy: 0.9714 - val_loss: 0.4061 - val_accuracy: 0.8945 -
lr: 1.0000e-05
Epoch 7/10
782/782 [=====] - ETA: 0s - loss: 0.0792 -
accuracy: 0.9724
Epoch 7: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0792 - accuracy: 0.9724 - val_loss: 0.4025 - val_accuracy: 0.8954 -
lr: 1.0000e-05
Epoch 8/10
782/782 [=====] - ETA: 0s - loss: 0.0845 -
accuracy: 0.9704
Epoch 8: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0845 - accuracy: 0.9704 - val_loss: 0.4034 - val_accuracy: 0.8950 -
lr: 1.0000e-05
Epoch 9/10
782/782 [=====] - ETA: 0s - loss: 0.0814 -
accuracy: 0.9713
Epoch 9: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 109ms/step - loss:
0.0814 - accuracy: 0.9713 - val_loss: 0.4027 - val_accuracy: 0.8949 -
lr: 1.0000e-05
Epoch 10/10
782/782 [=====] - ETA: 0s - loss: 0.0805 -
accuracy: 0.9721
Epoch 10: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0805 - accuracy: 0.9721 - val_loss: 0.4020 - val_accuracy: 0.8960 -
lr: 1.0000e-05

<keras.callbacks.History at 0x7fef0458a60>

```
model.fit(train_iterator, batch_size=batch_size, epochs=15, verbose=1, val  
idation_data=(X_test, y_test), callbacks =  
[checkpoint, reduce_lr, tensorboard_callback])
```

Epoch 1/15
782/782 [=====] - ETA: 0s - loss: 0.0834 -
accuracy: 0.9709
Epoch 1: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 109ms/step - loss:
0.0834 - accuracy: 0.9709 - val_loss: 0.4030 - val_accuracy: 0.8953 -
lr: 1.0000e-05
Epoch 2/15
782/782 [=====] - ETA: 0s - loss: 0.0796 -
accuracy: 0.9713
Epoch 2: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0796 - accuracy: 0.9713 - val_loss: 0.4020 - val_accuracy: 0.8952 -
lr: 1.0000e-05

Epoch 3/15
782/782 [=====] - ETA: 0s - loss: 0.0795 - accuracy: 0.9718
Epoch 3: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss: 0.0795 - accuracy: 0.9718 - val_loss: 0.4025 - val_accuracy: 0.8954 - lr: 1.0000e-05
Epoch 4/15
782/782 [=====] - ETA: 0s - loss: 0.0795 - accuracy: 0.9721
Epoch 4: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss: 0.0795 - accuracy: 0.9721 - val_loss: 0.4060 - val_accuracy: 0.8948 - lr: 1.0000e-05
Epoch 5/15
782/782 [=====] - ETA: 0s - loss: 0.0774 - accuracy: 0.9731
Epoch 5: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss: 0.0774 - accuracy: 0.9731 - val_loss: 0.4033 - val_accuracy: 0.8955 - lr: 1.0000e-05
Epoch 6/15
782/782 [=====] - ETA: 0s - loss: 0.0784 - accuracy: 0.9719
Epoch 6: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss: 0.0784 - accuracy: 0.9719 - val_loss: 0.4007 - val_accuracy: 0.8957 - lr: 1.0000e-05
Epoch 7/15
782/782 [=====] - ETA: 0s - loss: 0.0804 - accuracy: 0.9718
Epoch 7: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 108ms/step - loss: 0.0804 - accuracy: 0.9718 - val_loss: 0.4027 - val_accuracy: 0.8963 - lr: 1.0000e-05
Epoch 8/15
782/782 [=====] - ETA: 0s - loss: 0.0807 - accuracy: 0.9714
Epoch 8: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss: 0.0807 - accuracy: 0.9714 - val_loss: 0.4037 - val_accuracy: 0.8954 - lr: 1.0000e-05
Epoch 9/15
782/782 [=====] - ETA: 0s - loss: 0.0779 - accuracy: 0.9726
Epoch 9: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss: 0.0779 - accuracy: 0.9726 - val_loss: 0.4006 - val_accuracy: 0.8968 - lr: 1.0000e-05
Epoch 10/15

```

782/782 [=====] - ETA: 0s - loss: 0.0787 -
accuracy: 0.9719
Epoch 10: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 108ms/step - loss:
0.0787 - accuracy: 0.9719 - val_loss: 0.4000 - val_accuracy: 0.8967 -
lr: 1.0000e-05
Epoch 11/15
782/782 [=====] - ETA: 0s - loss: 0.0780 -
accuracy: 0.9720
Epoch 11: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0780 - accuracy: 0.9720 - val_loss: 0.4011 - val_accuracy: 0.8968 -
lr: 1.0000e-05
Epoch 12/15
782/782 [=====] - ETA: 0s - loss: 0.0760 -
accuracy: 0.9730
Epoch 12: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0760 - accuracy: 0.9730 - val_loss: 0.4006 - val_accuracy: 0.8968 -
lr: 1.0000e-05
Epoch 13/15
782/782 [=====] - ETA: 0s - loss: 0.0765 -
accuracy: 0.9732
Epoch 13: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
0.0765 - accuracy: 0.9732 - val_loss: 0.3999 - val_accuracy: 0.8974 -
lr: 1.0000e-05
Epoch 14/15
782/782 [=====] - ETA: 0s - loss: 0.0771 -
accuracy: 0.9719
Epoch 14: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0771 - accuracy: 0.9719 - val_loss: 0.4000 - val_accuracy: 0.8973 -
lr: 1.0000e-05
Epoch 15/15
782/782 [=====] - ETA: 0s - loss: 0.0755 -
accuracy: 0.9733
Epoch 15: val_accuracy did not improve from 0.89750
782/782 [=====] - 90s 115ms/step - loss:
0.0755 - accuracy: 0.9733 - val_loss: 0.4001 - val_accuracy: 0.8965 -
lr: 1.0000e-05

```

```
<keras.callbacks.History at 0x7fef5deef430>
```

```

data_generator = ImageDataGenerator(width_shift_range=0.2,
height_shift_range=0.2, rotation_range =
0.6, shear_range=0.1, zoom_range=0.4, horizontal_flip=True)
# prepare training iterator
train_iterator = data_generator.flow(X_train, y_train,
batch_size=batch_size)

```

```
# determine Loss function and Optimizer
model.compile(loss='sparse_categorical_crossentropy',
optimizer=Adam(learning_rate = 0.0001),metrics=['accuracy'])

reduce_lr = ReduceLR0nPlateau(monitor='val_accuracy',
factor=0.1,patience=3, min_lr=0.000001)

model.fit(train_iterator,batch_size=batch_size,epochs=15,verbose=1,validation_data=(X_test, y_test),callbacks =
[checkpoint,reduce_lr,tensorboard_callback])
```

```
Epoch 1/15
782/782 [=====] - ETA: 0s - loss: 0.0902 -
accuracy: 0.9688
Epoch 1: val_accuracy did not improve from 0.89750
782/782 [=====] - 123s 111ms/step - loss:
0.0902 - accuracy: 0.9688 - val_loss: 0.4174 - val_accuracy: 0.8957 -
lr: 1.0000e-04
Epoch 2/15
782/782 [=====] - ETA: 0s - loss: 0.0879 -
accuracy: 0.9690
Epoch 2: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 109ms/step - loss:
0.0879 - accuracy: 0.9690 - val_loss: 0.4029 - val_accuracy: 0.8948 -
lr: 1.0000e-04
Epoch 3/15
782/782 [=====] - ETA: 0s - loss: 0.0860 -
accuracy: 0.9688
Epoch 3: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
0.0860 - accuracy: 0.9688 - val_loss: 0.4063 - val_accuracy: 0.8942 -
lr: 1.0000e-04
Epoch 4/15
782/782 [=====] - ETA: 0s - loss: 0.0843 -
accuracy: 0.9701
Epoch 4: val_accuracy did not improve from 0.89750
782/782 [=====] - 86s 110ms/step - loss:
0.0843 - accuracy: 0.9701 - val_loss: 0.4213 - val_accuracy: 0.8921 -
lr: 1.0000e-04
Epoch 5/15
782/782 [=====] - ETA: 0s - loss: 0.0809 -
accuracy: 0.9711
Epoch 5: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 109ms/step - loss:
0.0809 - accuracy: 0.9711 - val_loss: 0.4009 - val_accuracy: 0.8966 -
lr: 1.0000e-05
Epoch 6/15
782/782 [=====] - ETA: 0s - loss: 0.0785 -
accuracy: 0.9732
Epoch 6: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
```

0.0785 - accuracy: 0.9732 - val_loss: 0.4021 - val_accuracy: 0.8962 -
lr: 1.0000e-05
Epoch 7/15
782/782 [=====] - ETA: 0s - loss: 0.0773 -
accuracy: 0.9728
Epoch 7: val_accuracy did not improve from 0.89750
782/782 [=====] - 86s 110ms/step - loss:
0.0773 - accuracy: 0.9728 - val_loss: 0.3999 - val_accuracy: 0.8965 -
lr: 1.0000e-05
Epoch 8/15
782/782 [=====] - ETA: 0s - loss: 0.0794 -
accuracy: 0.9717
Epoch 8: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
0.0794 - accuracy: 0.9717 - val_loss: 0.3976 - val_accuracy: 0.8969 -
lr: 1.0000e-05
Epoch 9/15
782/782 [=====] - ETA: 0s - loss: 0.0763 -
accuracy: 0.9730
Epoch 9: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 108ms/step - loss:
0.0763 - accuracy: 0.9730 - val_loss: 0.4007 - val_accuracy: 0.8963 -
lr: 1.0000e-05
Epoch 10/15
782/782 [=====] - ETA: 0s - loss: 0.0761 -
accuracy: 0.9730
Epoch 10: val_accuracy did not improve from 0.89750
782/782 [=====] - 86s 109ms/step - loss:
0.0761 - accuracy: 0.9730 - val_loss: 0.3970 - val_accuracy: 0.8971 -
lr: 1.0000e-05
Epoch 11/15
782/782 [=====] - ETA: 0s - loss: 0.0781 -
accuracy: 0.9727
Epoch 11: val_accuracy did not improve from 0.89750
782/782 [=====] - 87s 111ms/step - loss:
0.0781 - accuracy: 0.9727 - val_loss: 0.4001 - val_accuracy: 0.8968 -
lr: 1.0000e-05
Epoch 12/15
782/782 [=====] - ETA: 0s - loss: 0.0775 -
accuracy: 0.9731
Epoch 12: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
0.0775 - accuracy: 0.9731 - val_loss: 0.3988 - val_accuracy: 0.8968 -
lr: 1.0000e-05
Epoch 13/15
782/782 [=====] - ETA: 0s - loss: 0.0767 -
accuracy: 0.9726
Epoch 13: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 108ms/step - loss:
0.0767 - accuracy: 0.9726 - val_loss: 0.3995 - val_accuracy: 0.8971 -

```
lr: 1.0000e-05
Epoch 14/15
782/782 [=====] - ETA: 0s - loss: 0.0742 -
accuracy: 0.9731
Epoch 14: val_accuracy did not improve from 0.89750
782/782 [=====] - 84s 107ms/step - loss:
0.0742 - accuracy: 0.9731 - val_loss: 0.4019 - val_accuracy: 0.8964 -
lr: 1.0000e-06
Epoch 15/15
782/782 [=====] - ETA: 0s - loss: 0.0752 -
accuracy: 0.9737
Epoch 15: val_accuracy did not improve from 0.89750
782/782 [=====] - 86s 110ms/step - loss:
0.0752 - accuracy: 0.9737 - val_loss: 0.3990 - val_accuracy: 0.8966 -
lr: 1.0000e-06
```

```
<keras.callbacks.History at 0x7fef76051190>
```

```
data_generator = ImageDataGenerator(width_shift_range=0.3,
height_shift_range=0.3, rotation_range =
0.6, shear_range=0.3, zoom_range=0.4, horizontal_flip=True)
# prepare training iterator
train_iterator = data_generator.flow(X_train, y_train,
batch_size=batch_size)

reduce_lr = ReduceLRonPlateau(monitor='val_accuracy',
factor=0.1, patience=4, min_lr=0.000000001)

model.fit(train_iterator, batch_size=batch_size, epochs=3, verbose=1, vali
dation_data=(X_test, y_test), callbacks =
[checkpoint, reduce_lr, tensorboard_callback])
```

```
Epoch 1/3
782/782 [=====] - ETA: 0s - loss: 0.7970 -
accuracy: 0.7960
Epoch 1: val_accuracy did not improve from 0.89750
782/782 [=====] - 86s 110ms/step - loss:
0.7970 - accuracy: 0.7960 - val_loss: 0.5673 - val_accuracy: 0.8686 -
lr: 1.0000e-09
Epoch 2/3
782/782 [=====] - ETA: 0s - loss: 0.7903 -
accuracy: 0.7992
Epoch 2: val_accuracy did not improve from 0.89750
782/782 [=====] - 89s 114ms/step - loss:
0.7903 - accuracy: 0.7992 - val_loss: 0.5685 - val_accuracy: 0.8687 -
lr: 1.0000e-09
Epoch 3/3
782/782 [=====] - ETA: 0s - loss: 0.7877 -
accuracy: 0.7983
Epoch 3: val_accuracy did not improve from 0.89750
782/782 [=====] - 85s 109ms/step - loss:
```


0.7877 - accuracy: 0.7983 - val_loss: 0.5676 - val_accuracy: 0.8685 -
lr: 1.0000e-09

<keras.callbacks.History at 0x7feeed061910>

model.load_weights('/content/model_save/weights-15-0.8975.hdf5')

Test the model

score = model.evaluate(X_test, y_test, verbose=1)

print('Test loss:', score[0])

print('Test accuracy:', score[1])

313/313 [=====] - 6s 19ms/step - loss: 0.3961

- accuracy: 0.8975

Test loss: 0.3961257040500641

Test accuracy: 0.8974999785423279

Save the trained weights in to .h5 format

model.save_weights("DNST_model.h5")

print("Saved model to disk")

Saved model to disk