

## Transfer Learning Assignment

Download all the data in this rar\_file , it contains all the data required for the assignment. When you unrar the file you'll get the files in the following format:

path/to/the/image.tif,category

where the categories are numbered 0 to 15, in the following order:

There is a file named as 'labels\_final.csv' , it consists of two columns. First column is path which is the required path to the images and second is the class label.

*#the dataset that you are dealing with is quite large 3.7 GB and hence there are two methods to import the data to Colab*

*# Method 1- you can use gdown module to get the data directly from Google drive to Colab*

*# the syntax is as follows !gdown --id file\_id , for ex - running the below cell will import the rvl-cdip.rar dataset*

```
#!gdown --id 1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu
```

*# Method -2 you can also import the data using wget function*

```
#https://www.youtube.com/watch?v=BPUfVq7RaY8
```

*#unrar the file*

```
get_ipython().system_raw("unrar x rvl-cdip.rar")
```

## 2. On this image data, you have to train 3 types of models as given below You have to split the data into Train and Validation data.

```
'''#import all the required libraries
```

```
import tensorflow as tf
```

```
import os
```

```
import numpy as np
```

```
import pandas as pd
```

```
df=pd.read_csv('labels_final.csv',dtype=str)'''
```

```
{"type":"string"}
```

1. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data. or you can use this method also <https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

Note- In the reference notebook you were dealing with jpg images, in the given dataset you are dealing with tiff images. Imagedatagenerator works with both type of images. If you want to use custom data pipeline then you have to convert your tiff images to jpg images.

1. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.
2. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)
3. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

<https://www.appliedaicourse.com/lecture/11/applied-machine-learning-online-course/3426/code-example-cats-vs-dogs/8/module-8-neural-networks-computer-vision-and-deep-learning>

1. Do print model.summary() and draw model\_plots for each of the model.

### Import Section

```
import matplotlib.pyplot as plt # importing the libraries
import pandas as pd
import numpy as np
import os
import datetime
import seaborn as sns
import tensorflow as tf
import datetime, os
from tensorflow import keras
from keras_preprocessing.image import ImageDataGenerator
from keras.layers import Input, Lambda, Dense, Flatten , Conv2D,
MaxPool2D
from tensorflow.keras.callbacks import
ModelCheckpoint,EarlyStopping,LearningRateScheduler,ReduceLRonPlateau,
TensorBoard
from keras.models import Model
from keras.applications.vgg16 import VGG16,preprocess_input
from keras.preprocessing import image
from keras.callbacks import Callback,TensorBoard
from prettytable import PrettyTable
from prettytable import ALL as ALL
#!pip install -U --no-cache-dir gdown --pre
from tensorflow.keras.utils import plot_model
import warnings
warnings.filterwarnings('ignore')
```

### Lets Pull the Dataset

```
!pip install -q kaggle
from google.colab import files
files.upload()
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

```
!chmod 600 /root/.kaggle/kaggle.json
!mkdir dataset
! kaggle datasets download -d brahma0545/aaic-assignment-tl
! unzip /content/aaic-assignment-tl.zip -d dataset
```

```
dataframe=pd.read_csv("/content/dataset/labels_final.csv")
```

```
dataframe.head()
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

```
labels_dict={ 0 : "letter",1 : "form",2 : "email",3 : "handwritten",4 : "ad
vertisement",5 : "scientific report",6 : "scientific
publication",7 : "specification",8 : "file folder",9 : "news
article",10 : " budget",11 : "invoice",12 : "
presentation",13 : "questionnaire",14 : "resume",15: "memo"}
```

```
dataframe['label']=dataframe['label'].apply(lambda x:labels_dict[x])
dataframe.head()
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	handwritten
1	imagesl/l/x/t/lxt19d00/502213303.tif	handwritten
2	imagesx/x/e/d/xed05a00/2075325674.tif	email
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	handwritten
4	imagesq/q/z/k/qzk17e00/2031320195.tif	specification

## Lets Create Data Generator

```
datagenerator = ImageDataGenerator(rescale=1/255.,
validation_split=0.2) #image generator
```

```
print("-----TRAIN DATA-----") # train data
```

```
train_generator =
```

```
datagenerator.flow_from_dataframe(dataframe=dataframe,
directory="/content/dataset/data_final",
```

```
    x_col='path',
    y_col='label', # using
```

```
flow from data frame
```

```
    target_size=(256,256),
    class_mode='categorical',
    batch_size=32,
    subset='training',
    seed=7)
```

```
-----TRAIN DATA-----
```

```
Found 38400 validated image filenames belonging to 16 classes.
```

```

print("-----TEST DATA-----") # cross validation data
test_generator =
datagenerator.flow_from_dataframe(dataframe=dataframe,
directory="/content/dataset/data_final",
                                x_col='path',
                                y_col='label',
                                target_size=(256,256),
                                class_mode='categorical',
                                batch_size=32,
                                subset='validation',
                                seed=7)

```

-----TEST DATA-----

Found 9600 validated image filenames belonging to 16 classes.

```

result_dataset = pd.DataFrame(data = np.zeros((3,2)),index =
['Model_1', 'Model_2', 'Model_3'],columns =
['Model_Loss', 'Model_Accuracy'])
result_dataset

```

	Model_Loss	Model_Accuracy
Model_1	0.0	0.0
Model_2	0.0	0.0
Model_3	0.0	0.0

### Model-1

```

vgg_input = [256, 256] #pre trained vgg16 model
vgg_model = VGG16(input_shape=vgg_input + [3], weights='imagenet',
include_top=False)

```

```
vgg_model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168

block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0

=====  
Total params: 14,714,688  
Trainable params: 14,714,688  
Non-trainable params: 0

---

plot\_model(vgg\_model, show\_shapes = True)

input_1	input:	[(None, 256, 256, 3)]
InputLayer	output:	[(None, 256, 256, 3)]



block1_conv1	input:	(None, 256, 256, 3)
Conv2D	output:	(None, 256, 256, 64)



block1_conv2	input:	(None, 256, 256, 64)
Conv2D	output:	(None, 256, 256, 64)



block1_pool	input:	(None, 256, 256, 64)
MaxPooling2D	output:	(None, 128, 128, 64)



block2_conv1	input:	(None, 128, 128, 64)
Conv2D	output:	(None, 128, 128, 128)



block2_conv2	input:	(None, 128, 128, 128)
Conv2D	output:	(None, 128, 128, 128)



## Lets Build Model 1

- MODEL\_1(INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer )

```
# lets make vgg model layers non trainable
```

```
for layer in vgg_model.layers:
```

```
    layer.trainable = False
```

```
#Adding custom Layers
```

```
vgg_out = vgg_model.output
```

```
conv2d_1 =
```

```
Conv2D(filters=512, kernel_size=(3,3), padding="same", kernel_initializer  
='he_normal', activation="relu")(vgg_out)
```

```
maxpool_1 = MaxPool2D(2,2)(conv2d_1)
```

```
flatten_1 = Flatten()(maxpool_1)
```

```
dense_1 = Dense(256, activation="relu", kernel_initializer  
='he_normal')(flatten_1)
```

```
dense_2 = Dense(128, activation="relu", kernel_initializer  
='he_normal')(dense_1)
```

```
output = Dense(16, activation="softmax")(dense_2)
```

```
# creating the final model
```

```
model_1 = Model(inputs = vgg_model.input, outputs = output)
```

```
# compile the model
```

```
model_1.compile(loss = "categorical_crossentropy", optimizer = 'Adam',  
metrics=["accuracy"])
```

```
# summary of the model_1
```

```
model_1.summary()
```

Model: "model\_5"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080

block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_9 (Conv2D)	(None, 8, 8, 512)	2359808
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_5 (Flatten)	(None, 8192)	0
dense_7 (Dense)	(None, 256)	2097408
dense_8 (Dense)	(None, 128)	32896
dense_9 (Dense)	(None, 16)	2064

```
=====
Total params: 19,206,864
Trainable params: 4,492,176
Non-trainable params: 14,714,688
```

---

```
#lets plot model 1
plot_model(model_1,to_file='model_1.png',show_shapes = True)
```



input_1	input:	[(None, 256, 256, 3)]
InputLayer	output:	[(None, 256, 256, 3)]



block1_conv1	input:	(None, 256, 256, 3)
Conv2D	output:	(None, 256, 256, 64)



block1_conv2	input:	(None, 256, 256, 64)
Conv2D	output:	(None, 256, 256, 64)



block1_pool	input:	(None, 256, 256, 64)
MaxPooling2D	output:	(None, 128, 128, 64)



block2_conv1	input:	(None, 128, 128, 64)
Conv2D	output:	(None, 128, 128, 128)



block2_conv2	input:	(None, 128, 128, 128)
Conv2D	output:	(None, 128, 128, 128)



```

# HyperParameters
numberofepochs = 5
batch_size = 32 #note should be same as imagedatagenerator

filepath="model_1_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath,
monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.01,
patience=3, verbose=1)

# Load the TensorBoard notebook extension
%load_ext tensorboard
log_dir = os.path.join("logs", 'fits',
datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback =
tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,write_
graph=True)
%reload_ext tensorboard

reduce_lr = ReduceLRonPlateau(monitor='val_loss',
factor=0.2,patience=5, min_lr=0.000001)

#fitting the model_1
model_1.fit_generator(train_generator,epochs=numberofepochs,validation
_data=test_generator,callbacks=[checkpoint,earlystop,reduce_lr,tensorb
oard_callback])

Epoch 1/5
1200/1200 [=====] - ETA: 0s - loss: 1.2606 -
accuracy: 0.6137
Epoch 1: val_accuracy improved from -inf to 0.68667, saving model to
model_1_save/weights-01-0.6867.hdf5
1200/1200 [=====] - 321s 255ms/step - loss:
1.2606 - accuracy: 0.6137 - val_loss: 1.0319 - val_accuracy: 0.6867 -
lr: 0.0010
Epoch 2/5
1200/1200 [=====] - ETA: 0s - loss: 0.8603 -
accuracy: 0.7344
Epoch 2: val_accuracy improved from 0.68667 to 0.70781, saving model
to model_1_save/weights-02-0.7078.hdf5
1200/1200 [=====] - 296s 246ms/step - loss:
0.8603 - accuracy: 0.7344 - val_loss: 0.9810 - val_accuracy: 0.7078 -
lr: 0.0010
Epoch 3/5
1200/1200 [=====] - ETA: 0s - loss: 0.7237 -
accuracy: 0.7761
Epoch 3: val_accuracy improved from 0.70781 to 0.75104, saving model
to model_1_save/weights-03-0.7510.hdf5
1200/1200 [=====] - 296s 247ms/step - loss:
0.7237 - accuracy: 0.7761 - val_loss: 0.8553 - val_accuracy: 0.7510 -

```

```

lr: 0.0010
Epoch 4/5
1200/1200 [=====] - ETA: 0s - loss: 0.6082 -
accuracy: 0.8094
Epoch 4: val_accuracy improved from 0.75104 to 0.76448, saving model
to model_1_save/weights-04-0.7645.hdf5
1200/1200 [=====] - 296s 247ms/step - loss:
0.6082 - accuracy: 0.8094 - val_loss: 0.8355 - val_accuracy: 0.7645 -
lr: 0.0010
Epoch 5/5
1200/1200 [=====] - ETA: 0s - loss: 0.5162 -
accuracy: 0.8361
Epoch 5: val_accuracy improved from 0.76448 to 0.76875, saving model
to model_1_save/weights-05-0.7688.hdf5
1200/1200 [=====] - 297s 247ms/step - loss:
0.5162 - accuracy: 0.8361 - val_loss: 0.8470 - val_accuracy: 0.7688 -
lr: 0.0010

```

<keras.callbacks.History at 0x7ff7f4e94370>

```

model1_score = model_1.evaluate(test_generator, verbose=1)
print('Test loss:', model1_score[0])
print('Test accuracy:', model1_score[1])
result_dataset['Model_Loss']['Model_1'] = model1_score[0]
result_dataset['Model_Accuracy']['Model_1'] = model1_score[1]

300/300 [=====] - 58s 189ms/step - loss:
0.8470 - accuracy: 0.7688
Test loss: 0.8469867706298828
Test accuracy: 0.768750011920929

```

result\_dataset

	Model_Loss	Model_Accuracy
Model_1	0.846987	0.76875
Model_2	0.000000	0.000000
Model_3	0.000000	0.000000

%tensorboard --logdir logs

Output hidden; open in <https://colab.research.google.com> to view.

## Model-2

### Lets Build Model 2

MODEL\_2 ( INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer )

#model\_2

```

for layer in vgg_model.layers:
    layer.trainable = False

```

```

#Adding custom Layers
vgg_model_output = vgg_model.output
conv2d_1 =
Conv2D(filters=256,kernel_size=8 ,strides=1,kernel_initializer =
'he_normal',activation="relu")(vgg_model_output)
conv2d_2 =
Conv2D(filters=128,kernel_size=1 ,strides=1,kernel_initializer =
'he_normal',activation="relu")(conv2d_1)
flatten_1 = Flatten()(conv2d_2)
# creating the final model
output= Dense(16, activation="softmax")(flatten_1)
model_2 = Model(inputs = vgg_model.input, outputs = output)
# compile the model
model_2.compile(loss="categorical_crossentropy",optimizer =
'Adam',metrics=['accuracy'])

# summary of the model_2
model_2.summary()

```

Model: "model\_2"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808

block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_4 (Conv2D)	(None, 1, 1, 256)	8388864
conv2d_5 (Conv2D)	(None, 1, 1, 128)	32896
flatten_2 (Flatten)	(None, 128)	0
dense_2 (Dense)	(None, 16)	2064

```
=====
Total params: 23,138,512
Trainable params: 8,423,824
Non-trainable params: 14,714,688
```

---

```
plot_model(model_2,to_file = 'model_2.png',show_shapes = True)
```

input_1	input:	[(None, 256, 256, 3)]
InputLayer	output:	[(None, 256, 256, 3)]



block1_conv1	input:	(None, 256, 256, 3)
Conv2D	output:	(None, 256, 256, 64)



block1_conv2	input:	(None, 256, 256, 64)
Conv2D	output:	(None, 256, 256, 64)



block1_pool	input:	(None, 256, 256, 64)
MaxPooling2D	output:	(None, 128, 128, 64)



block2_conv1	input:	(None, 128, 128, 64)
Conv2D	output:	(None, 128, 128, 128)



block2_conv2	input:	(None, 128, 128, 128)
Conv2D	output:	(None, 128, 128, 128)



```
filepath="model_2_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath,
monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')
```

```
#fitting the model_2
```

```
model_2.fit_generator(train_generator, epochs=numberofepochs, validation
_data=test_generator, callbacks=[checkpoint, earlystop, reduce_lr, tensorb
oard_callback])
```

```
Epoch 1/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 1.2592 -
accuracy: 0.6249
```

```
Epoch 1: val_accuracy improved from -inf to 0.68313, saving model to
model_2_save/weights-01-0.6831.hdf5
```

```
1200/1200 [=====] - 325s 260ms/step - loss:
1.2592 - accuracy: 0.6249 - val_loss: 1.0398 - val_accuracy: 0.6831 -
lr: 0.0010
```

```
Epoch 2/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 0.8692 -
accuracy: 0.7333
```

```
Epoch 2: val_accuracy improved from 0.68313 to 0.72260, saving model
to model_2_save/weights-02-0.7226.hdf5
```

```
1200/1200 [=====] - 314s 261ms/step - loss:
0.8692 - accuracy: 0.7333 - val_loss: 0.9323 - val_accuracy: 0.7226 -
lr: 0.0010
```

```
Epoch 3/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 0.7271 -
accuracy: 0.7765
```

```
Epoch 3: val_accuracy improved from 0.72260 to 0.74406, saving model
to model_2_save/weights-03-0.7441.hdf5
```

```
1200/1200 [=====] - 313s 260ms/step - loss:
0.7271 - accuracy: 0.7765 - val_loss: 0.8823 - val_accuracy: 0.7441 -
lr: 0.0010
```

```
Epoch 4/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 0.6209 -
accuracy: 0.8063
```

```
Epoch 4: val_accuracy improved from 0.74406 to 0.74542, saving model
to model_2_save/weights-04-0.7454.hdf5
```

```
1200/1200 [=====] - 310s 258ms/step - loss:
0.6209 - accuracy: 0.8063 - val_loss: 0.9209 - val_accuracy: 0.7454 -
lr: 0.0010
```

```
Epoch 5/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 0.5350 -
accuracy: 0.8292
```

```
Epoch 5: val_accuracy did not improve from 0.74542
```

```
1200/1200 [=====] - 309s 258ms/step - loss:
0.5350 - accuracy: 0.8292 - val_loss: 0.9859 - val_accuracy: 0.7440 -
lr: 0.0010
```

```
<keras.callbacks.History at 0x7f176ce63df0>
```

```

model2_score = model_2.evaluate(test_generator, verbose=1)
print('Test loss:', model2_score[0])
print('Test accuracy:', model2_score[1])
result_dataset['Model_Loss']['Model_2'] = model2_score[0]
result_dataset['Model_Accuracy']['Model_2'] = model2_score[1]

300/300 [=====] - 56s 187ms/step - loss:
0.9209 - accuracy: 0.7454
Test loss: 0.9209381341934204
Test accuracy: 0.7454166412353516

```

result\_dataset

	Model_Loss	Model_Accuracy
Model_1	0.846987	0.768750
Model_2	0.920938	0.745417
Model_3	0.000000	0.000000

```
%tensorboard --logdir logs
```

Output hidden; open in <https://colab.research.google.com> to view.

```
tf.keras.backend.clear_session
```

```
<function keras.backend.clear_session()>
```

## Model-3

### Lets Build Model-3

- 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
# lets set last 6 layers trainable
```

```
for layer in vgg_model.layers[-6:]: # training last 6 layers of vgg16
    layer.trainable = True
    print("Layer '%s' is trainable" % layer.name)
```

```
Layer 'block4_conv3' is trainable
```

```
Layer 'block4_pool' is trainable
```

```
Layer 'block5_conv1' is trainable
```

```
Layer 'block5_conv2' is trainable
```

```
Layer 'block5_conv3' is trainable
```

```
Layer 'block5_pool' is trainable
```

```
#model_3
```

```
#Adding custom Layers
```

```
vgg_output = vgg_model.output
```

```
conv2d_1 =
```

```
Conv2D(filters=256, kernel_size=8, strides=1, kernel_initializer =
'he_normal', activation="relu")(vgg_output)
```

```
conv2d_2 =
```



```

Conv2D(filters=128,kernel_size=1 ,strides=1,kernel_initializer =
'he_normal',activation="relu")(conv2d_1)
flatten_1 = Flatten()(conv2d_2)
# creating the final model
output = Dense(16, activation="softmax")(flatten_1)
model_3 = Model(inputs = vgg_model.input, outputs = output)
# compile the model
model_3.compile(loss="categorical_crossentropy",optimizer =
'Adam',metrics=['accuracy'])

```

```
model_3.summary()
```

```
Model: "model_3"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808

block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_6 (Conv2D)	(None, 1, 1, 256)	8388864
conv2d_7 (Conv2D)	(None, 1, 1, 128)	32896
flatten_3 (Flatten)	(None, 128)	0
dense_3 (Dense)	(None, 16)	2064

```
=====
Total params: 23,138,512
Trainable params: 17,863,056
Non-trainable params: 5,275,456
```

---

```
plot_model(model_3,to_file = 'model_3.png',show_shapes = True)
```

input_1	input:	[(None, 256, 256, 3)]
InputLayer	output:	[(None, 256, 256, 3)]



block1_conv1	input:	(None, 256, 256, 3)
Conv2D	output:	(None, 256, 256, 64)



block1_conv2	input:	(None, 256, 256, 64)
Conv2D	output:	(None, 256, 256, 64)



block1_pool	input:	(None, 256, 256, 64)
MaxPooling2D	output:	(None, 128, 128, 64)



block2_conv1	input:	(None, 128, 128, 64)
Conv2D	output:	(None, 128, 128, 128)



block2_conv2	input:	(None, 128, 128, 128)
Conv2D	output:	(None, 128, 128, 128)



```
filepath="model_3_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath,
monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')
```

```
#fitting the model_3
```

```
model_3.fit_generator(train_generator, epochs=numberofepochs, validation
_data=test_generator, callbacks=[checkpoint, earlystop, reduce_lr, tensorb
oard_callback])
```

```
Epoch 1/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 2.7925 -
accuracy: 0.0608
```

```
Epoch 1: val_accuracy improved from -inf to 0.05844, saving model to
model_3_save/weights-01-0.0584.hdf5
```

```
1200/1200 [=====] - 411s 341ms/step - loss:
2.7925 - accuracy: 0.0608 - val_loss: 2.7729 - val_accuracy: 0.0584 -
lr: 0.0010
```

```
Epoch 2/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 2.7729 -
accuracy: 0.0607
```

```
Epoch 2: val_accuracy did not improve from 0.05844
```

```
1200/1200 [=====] - 355s 295ms/step - loss:
2.7729 - accuracy: 0.0607 - val_loss: 2.7730 - val_accuracy: 0.0584 -
lr: 0.0010
```

```
Epoch 3/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 2.7728 -
accuracy: 0.0594
```

```
Epoch 3: val_accuracy did not improve from 0.05844
```

```
1200/1200 [=====] - 355s 296ms/step - loss:
2.7728 - accuracy: 0.0594 - val_loss: 2.7732 - val_accuracy: 0.0584 -
lr: 0.0010
```

```
Epoch 4/5
```

```
1200/1200 [=====] - ETA: 0s - loss: 2.7728 -
accuracy: 0.0620
```

```
Epoch 4: val_accuracy improved from 0.05844 to 0.06115, saving model
to model_3_save/weights-04-0.0611.hdf5
```

```
1200/1200 [=====] - 355s 296ms/step - loss:
2.7728 - accuracy: 0.0620 - val_loss: 2.7730 - val_accuracy: 0.0611 -
lr: 0.0010
```

```
Epoch 4: early stopping
```

```
<keras.callbacks.History at 0x7f16e66519a0>
```

```
model3_score = model_3.evaluate(test_generator, verbose=1)
```

```
print('Test loss:', model3_score[0])
```

```
print('Test accuracy:', model3_score[1])
```

```
result_dataset['Model_Loss']['Model_3'] = model3_score[0]
```

```
result_dataset['Model_Accuracy']['Model_3'] = model3_score[1]
```

```
300/300 [=====] - 58s 190ms/step - loss:
2.7730 - accuracy: 0.0611
```

```
Test loss: 2.7729671001434326
Test accuracy: 0.0611458346247673
```

```
result_dataset
```

```
      Model_Loss  Model_Accuracy
Model_1    0.846987      0.768750
Model_2    0.920938      0.745417
Model_3    2.772967      0.061146
```

```
%tensorboard --logdir logs
```

Output hidden; open in <https://colab.research.google.com> to view.

## Observations

Please write your observations or a brief summary of the results that you get after performing transfer learning with reference to model1, model2 and model3

```
from tabulate import tabulate
print(tabulate(result_dataset, headers='keys', tablefmt='psql'))
```

```
+-----+-----+-----+
|      | Model_Loss | Model_Accuracy |
+-----+-----+-----+
| Model_1 |    0.846987 |    0.76875    |
| Model_2 |    0.920938 |    0.745417   |
| Model_3 |    2.77297  |    0.0611458  |
+-----+-----+-----+
```

## Observations

### model 1:

- We can see that transfer learning using top layers have very good accuracy as the weights are well converged imagenet weights. And accuracy is around 77% percentage for 5 epochs, if we increase the epochs we will get very good accuracy in few epochs.and we have around 19 million weights for this model.and train time it is around 300 seconds.

### ##model 2:

- here we are using transfer learning using top layers with conv layers instead of dense layers to retain the spatial information .And accuracy is around 74.5% percentage for 5 epochs, if we increase the epochs we will get very good accuracy in few epochs.and we have around 23 million parameters for this model.which increased when compared to model 1.and train time increased to 325 seconds.

### ##model 3:

- here we are using transfer learning using top layers and last six layers trainable with conv layers instead of dense layers to retain the spatial information .And accuracy is around 6.11% percentage for 5 epochs as we are re training six layers from vgg 16 it takes time to converge for our dataset .and trainable parameters increased to 23.1 million parameters for this model.which increased when compared to model 2.and train time increased to 355 seconds.